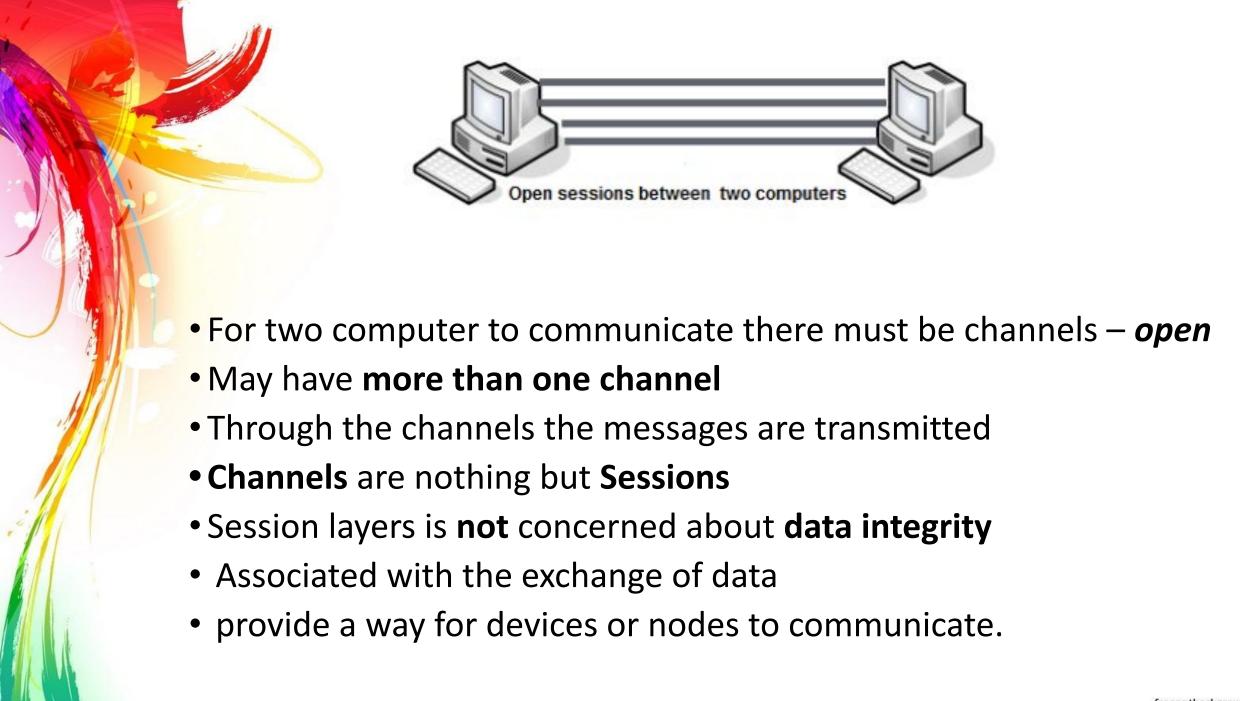




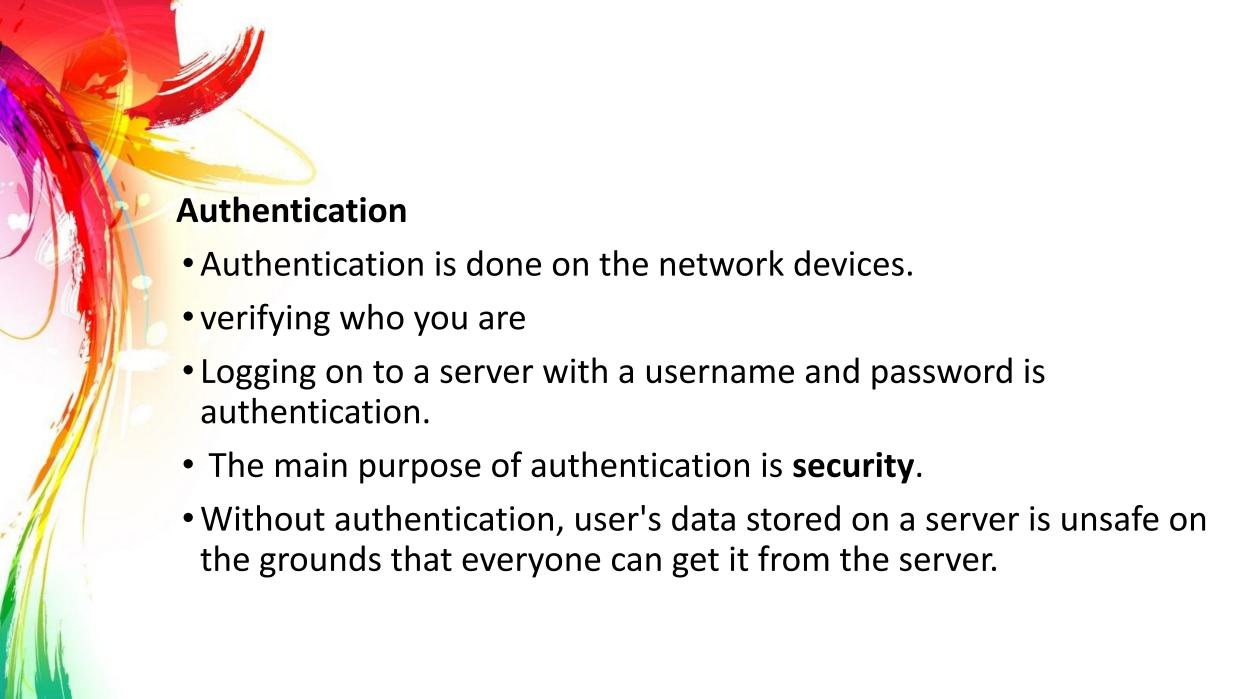
Session Layer

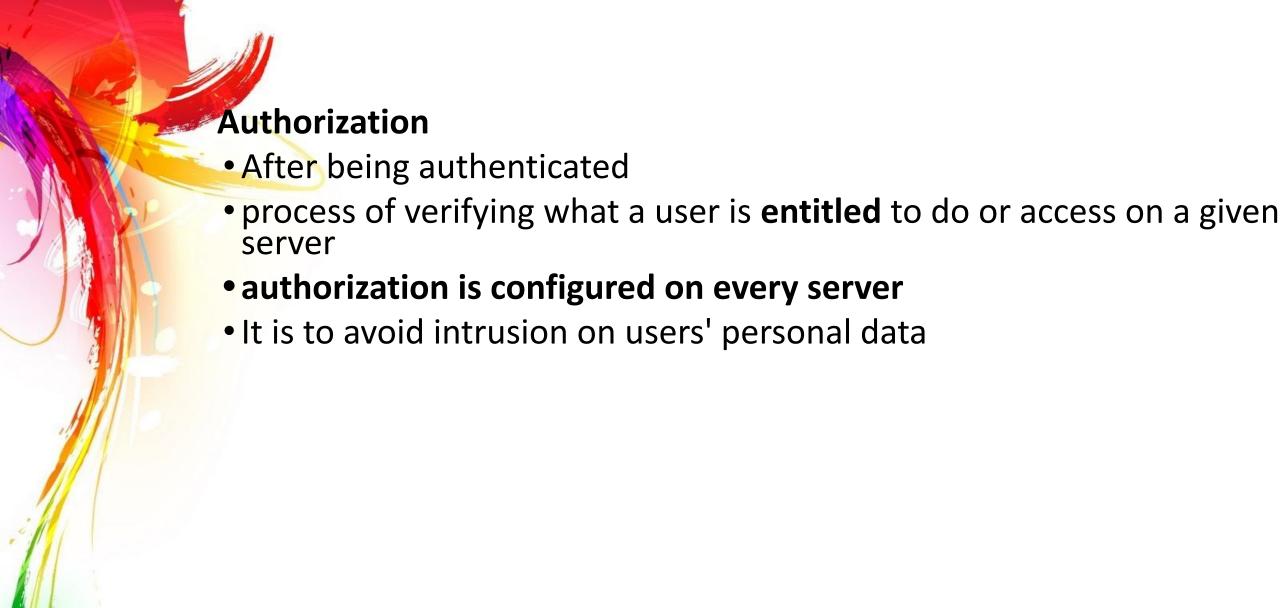
Session Layer

- 5th layer of OSI Reference Model
- Responsible for
 - ✓ Opening session
 - ✓ Managing session
 - ✓ Terminating session









Dialog control

- A successful communication is the one that is well organized and sorted out.
- a third party in a communicative context may effectively lead to a successful communication by regulating things.
- This third party in computer network is called 'Dialog control' and 'Dialog separation'.

☐ dialog control

- it determines whose turn it is to transfer data in a session
- In a given open session, a device plays dual roles, which is requesting services and replying with services.

 Dialog control determines which role they are playing at any given moment.

dialog control is critically significant for data transmission.

provides three different modes or ways of communication, which

are

✓ simplex

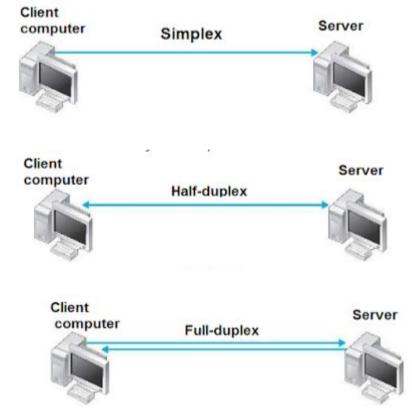
eg: Radio

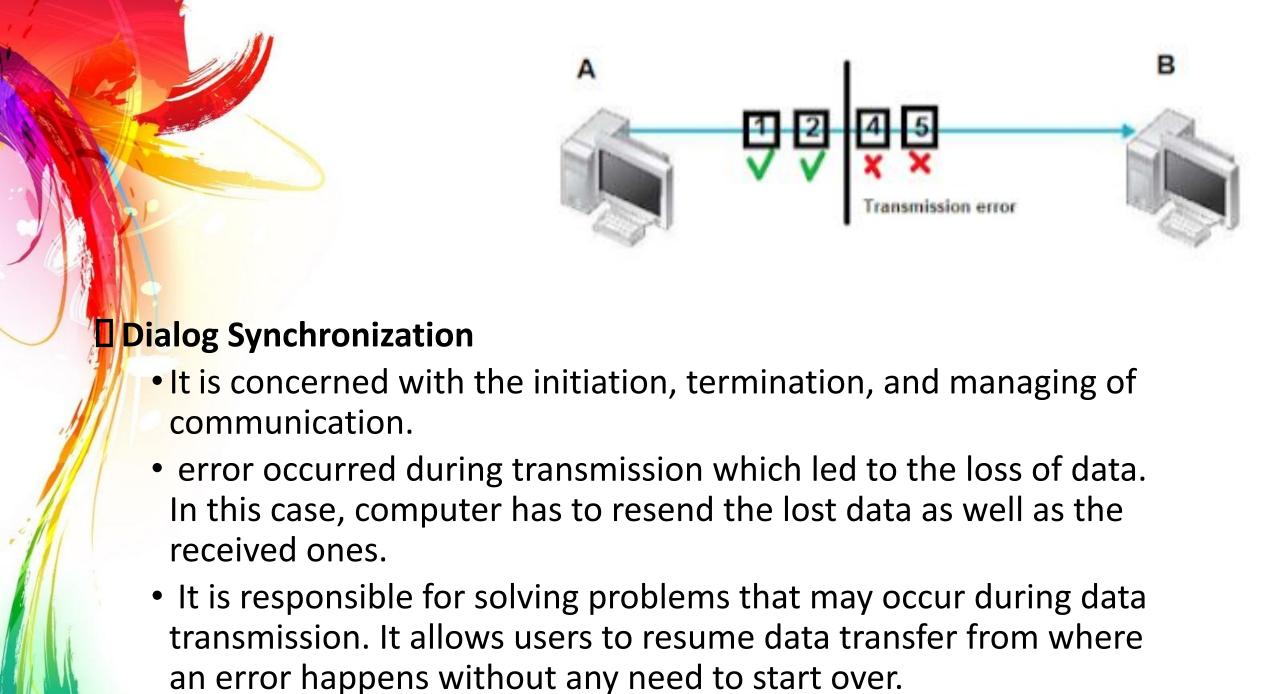
✓ half-duplex

eg: Walkie talkie

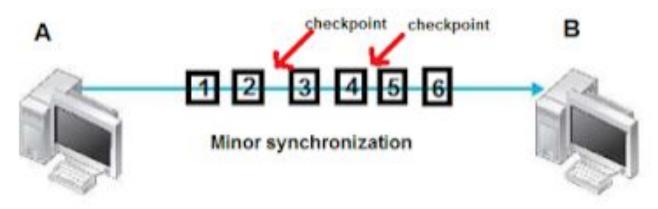
✓ full-duplex

eg: skype





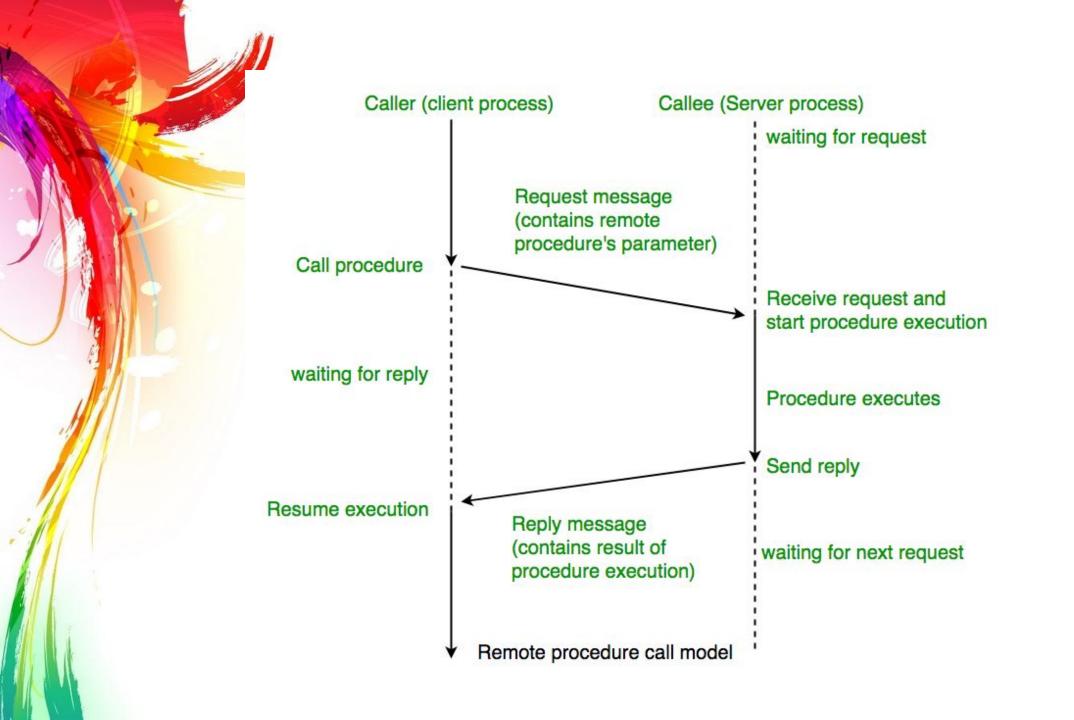
- when an error occurs during transmission, the sender does not need to start sending data from the beginning, it rather start from where it breaks. This is achieved through synchronization.
- If the session is closed by mistake, the dialog separation has the ability to re-establish the connection and start from where it leaves off.
- During the transmission of data, dialog separation allows a process to add checkpoints to each set of data.

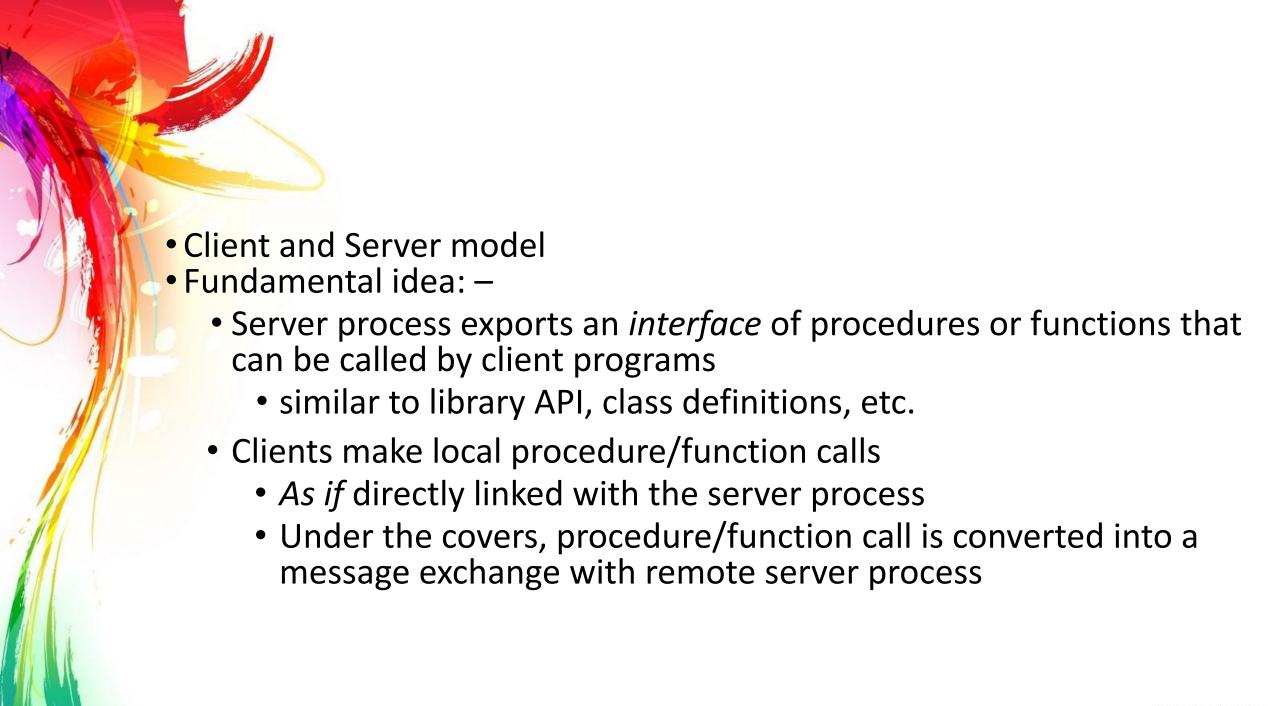


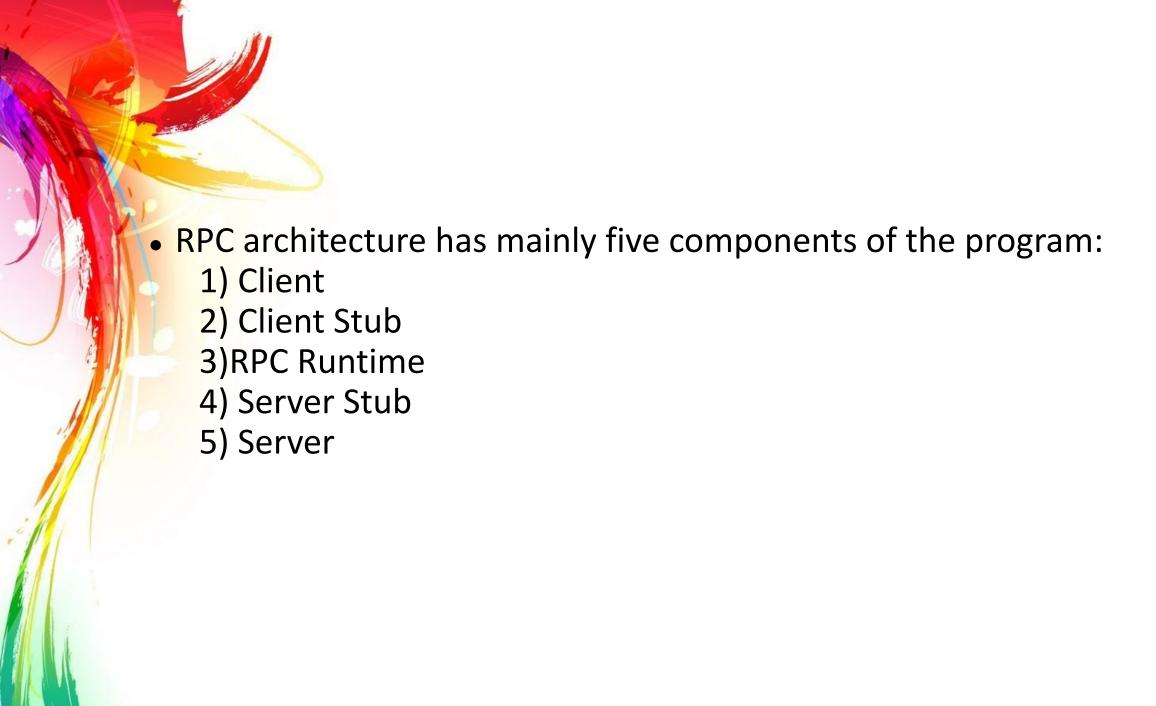
- each piece of data is assigned a *sequence number*
- when an error occurs, the receiver can re-synchronize the state of the session to a previous synchronization point.

Remote Procedure Call

- Remote Procedure Call (RPC) provides a different paradigm for accessing network services.
- Making a call to a function to request for a service, on another machine without having to know the networking details
- Session-layer services are commonly used in application environments that make use of Remote Procedure Calls (RPCs)
- Instead of accessing remote services by sending and receiving messages, a client invokes services by making a local procedure call.
- The local procedure hides the details of the network communication.
- Message based communication system

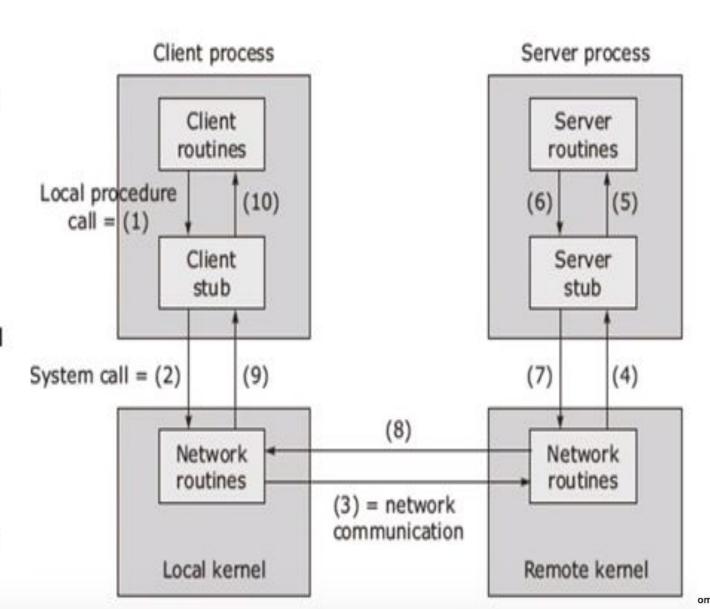






RPC Execution

- Client calls a local procedure on the client stub
- The client stub acts as a proxy and marshalls the call and the args.
- The client stub send this to the remote system (via TCP/UDP)
- The server stub unmarshalls the call and args from the client
- The server stub calls the actual procedure on the server
- The server stub marshalls the reply and sends it back to the client



Client

 A process, such as a program or task, that requests a service provided by another program. The client process uses the requested service without having to "deal" with many working details about the other program or the service.

Server

• A process, such as a program or task, that responds to requests from a client.

Client Stub

 Module within a client application containing all of the functions necessary for the client to make remote procedure calls using the model of a traditional function call in a standalone application. The client stub is responsible for invoking the marshalling engine and some of the RPC application programming interfaces (APIs).

Server Stub

 Module within a server application or service that contains all of the functions necessary for the server to handle remote requests using local procedure calls.

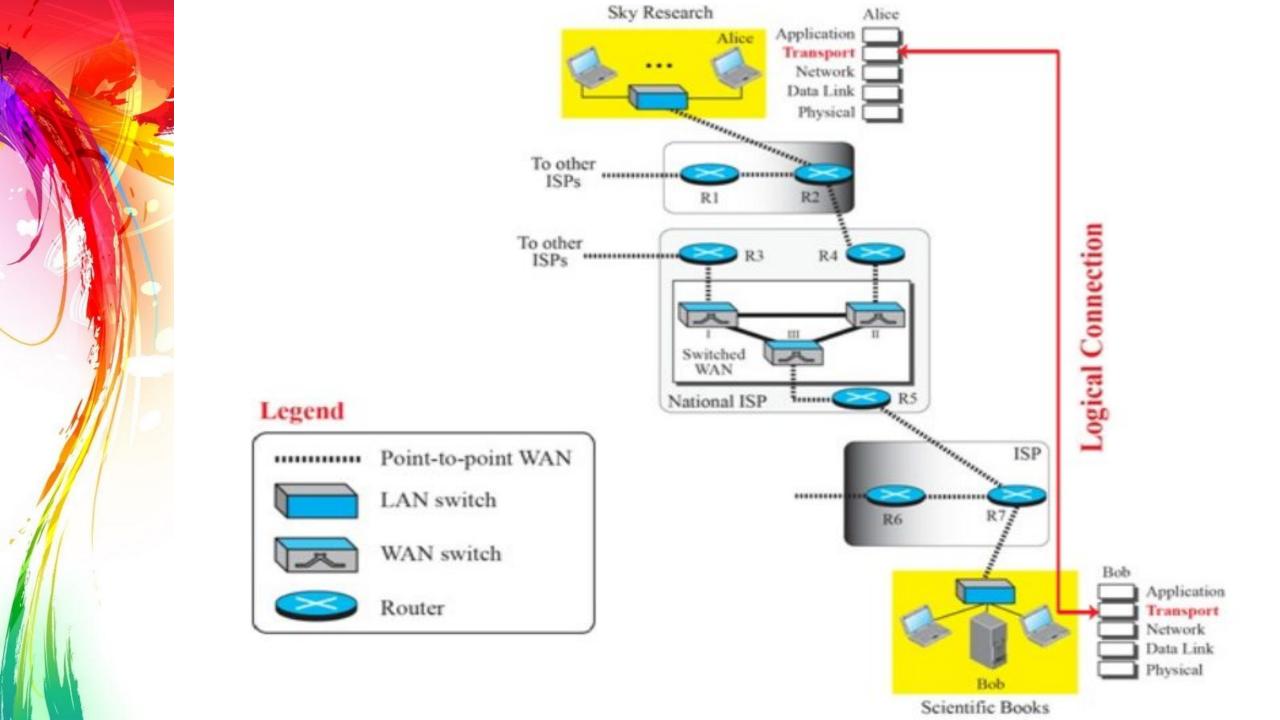
- RPC is supported by no of network protocols TCP, UDP
- Marshalling Arguments
 - Marshalling is the packing of function parameters into a message packet
 - the RPC stubs call type-specific functions to marshal or unmarshal the parameters of an RPC
 - Client stub marshals the arguments into a message
 - Server stub unmarshals the arguments and uses them to invoke the service function
 - on return:
 - the server stub marshals return values
 - the client stub unmarshals return values, and returns to the client program

 freepptbackground.com



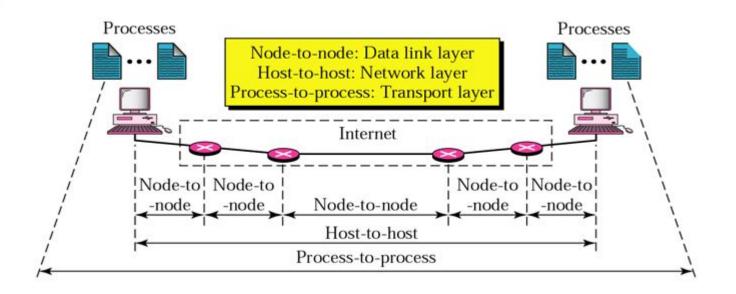
Transport Layer

- The transport layer is located between the application layer and the network layer in TCP reference Model, and between session and network layer in OSI reference Model.
- It provides end to end and process-to-process communication between two application layers, one at the local host and the other at the remote host.
- Communication is provided using a *logical connection*.
- The transport layer is responsible for *providing services to the* application layer/ session layer; it receives services from the network layer.



Network Layer Vs Transport Layer

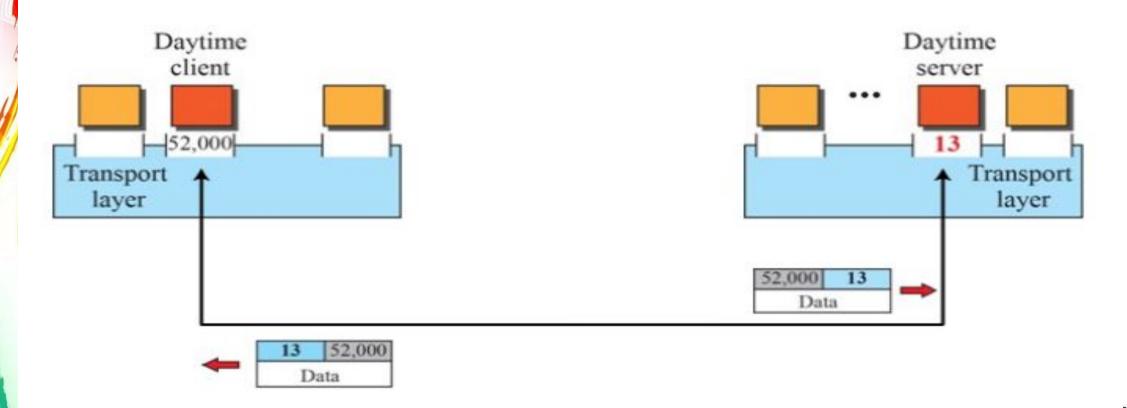
 Network layer only does host to host delivery, later the transport layer takes over to assign to process (port)

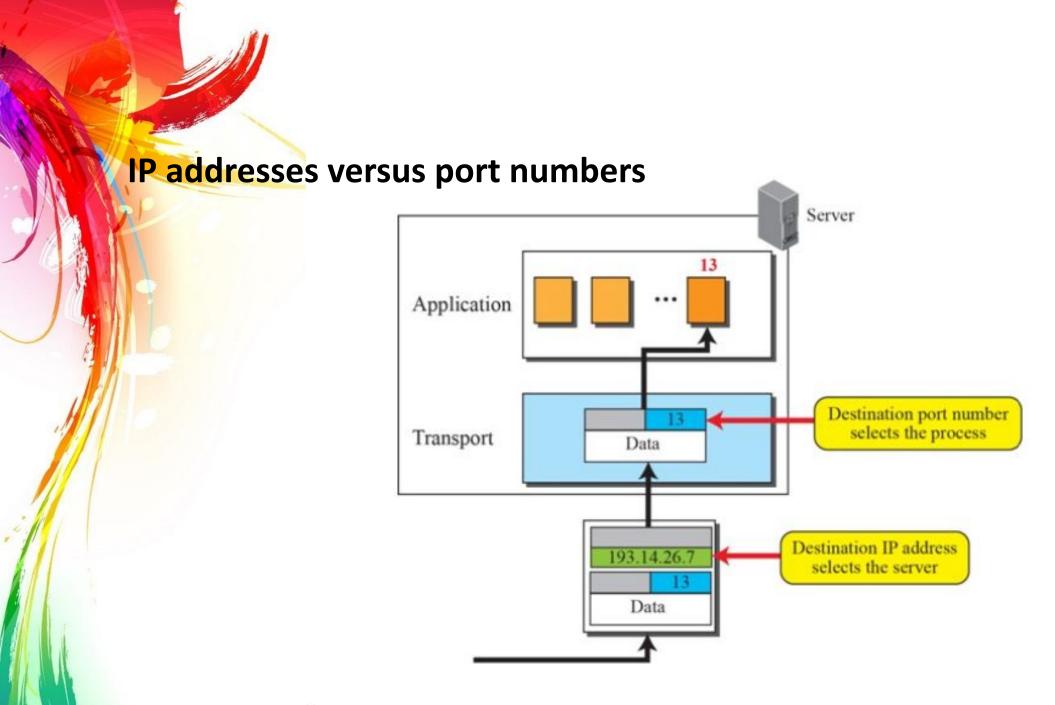


• The transport layer is responsible for **process-to-process** delivery

Port numbers

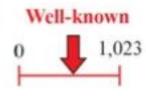
- Both process will have same name
- To find process find port no

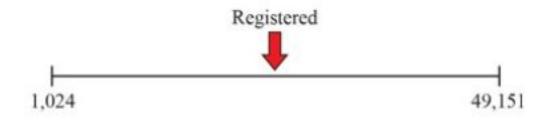


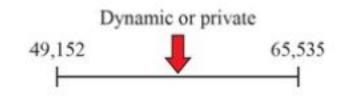


freepptbackground.com

- ICANN (Internet Corporation for Assigned Names and Numbers)ranges
- Port nos use 16 bits (2¹⁶)









In UNIX, the well-known ports are stored in a file called /etc/services. We can use the *grep* utility to extract the line corresponding to the desired application.

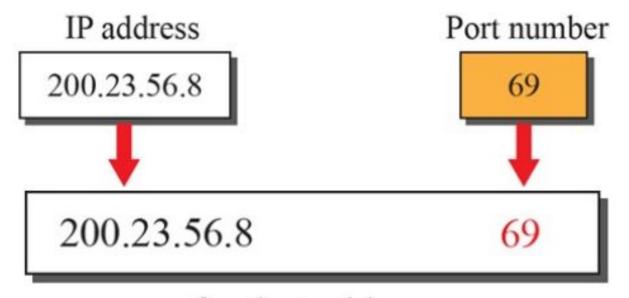
```
$grep tftp/etc/services
tftp 69/tcp
tftp 69/udp
```

\$grep snmp/etc/services snmp161/tcp#Simple Net Mgmt Proto snmp161/udp#Simple Net Mgmt Proto snmptrap162/udp#Traps for SNMP

TFTP uses 69
SNMP uses 161 and 162



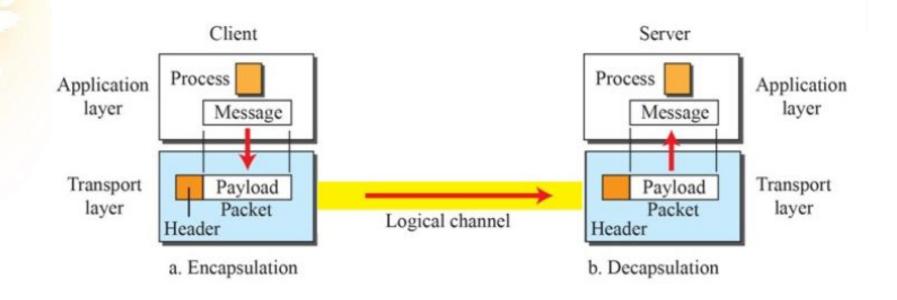
Transport layer protocol need socket address



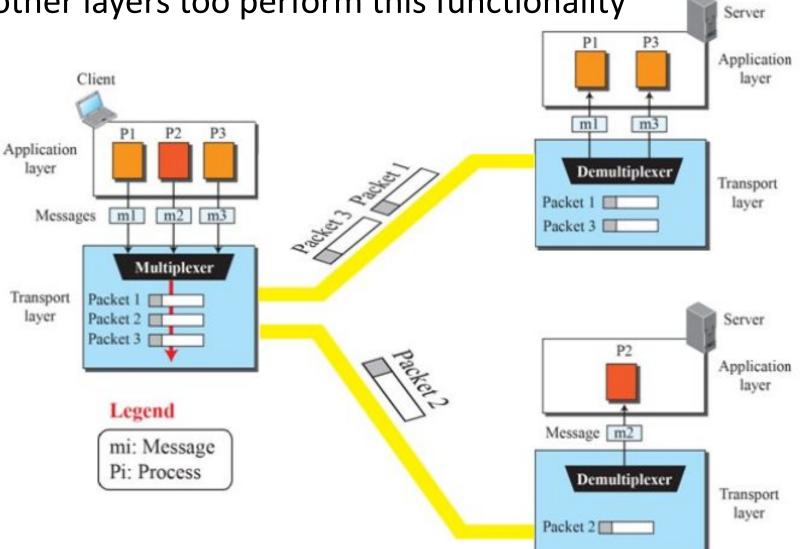
Functionality of Transport Layer

Encapsulation and encapsulation

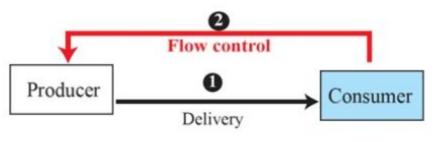
- User datagram/ segment are encapsulated
- other layers too perform this functionality



Multiplexing and demultiplexing other layers too perform this functionality



Flow Control



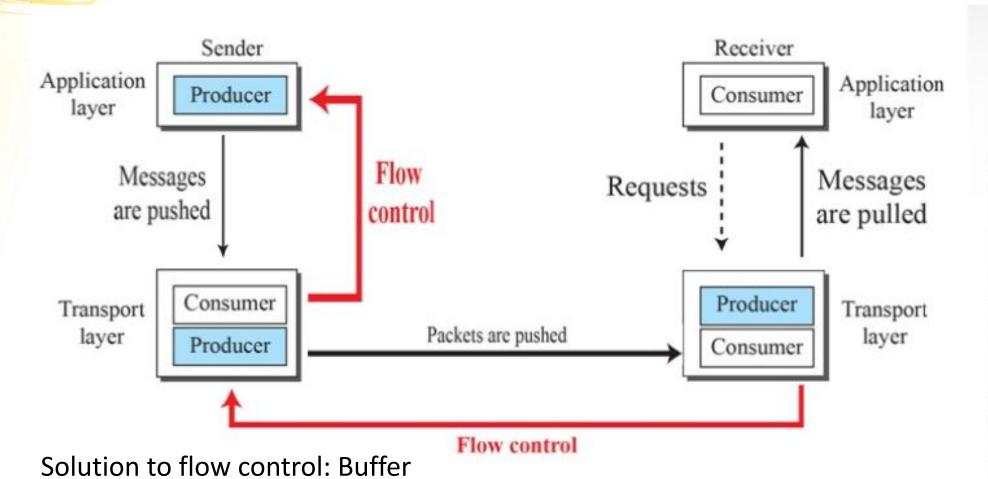
a. Pushing



b. Pulling

Flow control at the transport layer

- The production and consumption rate must be balanced
- If consumer or producer is overloaded, it has to signal to the other to stop



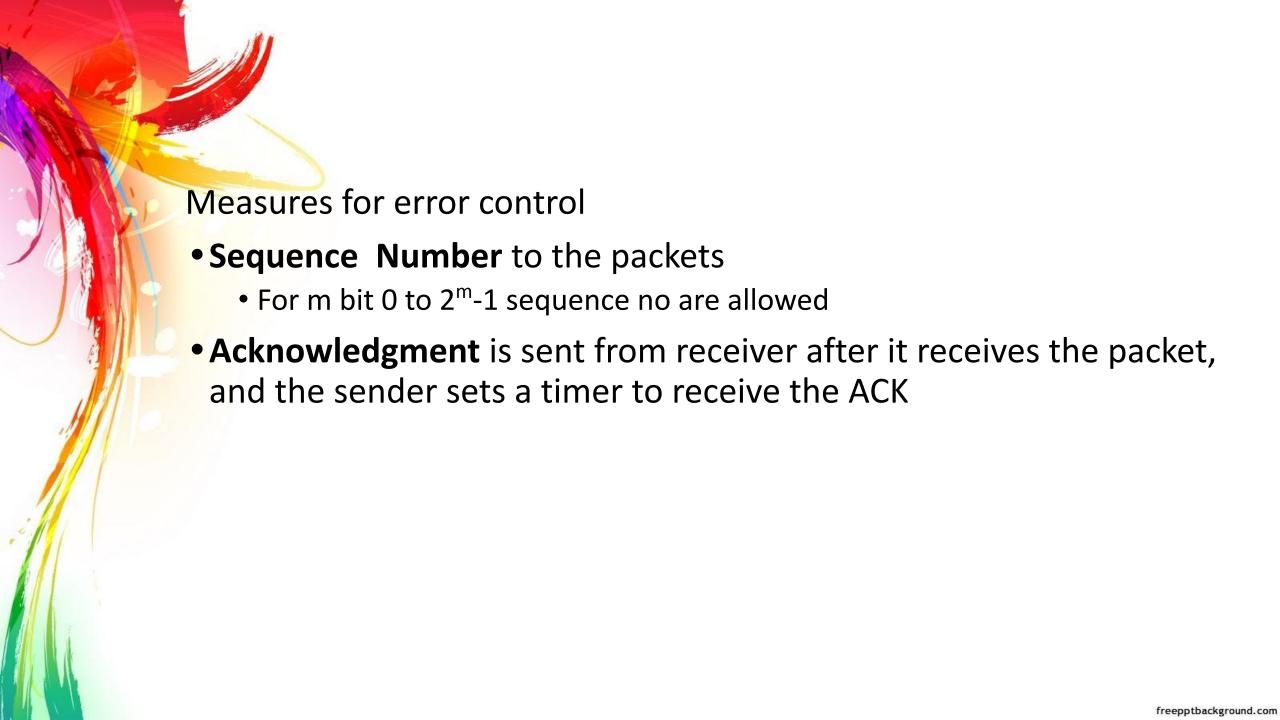
• When buffer is full the receiver informs the sender to stop



Error Control

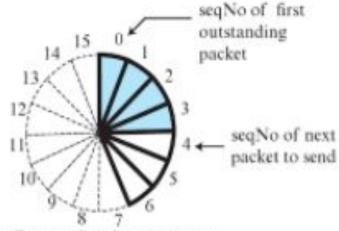
- 1. Detecting and discarding corrupted packets.
- 2. Keeping track of lost and discarded packets and resending them.
- 3. Recognizing duplicate packets and discarding them.
- 4. Buffering out-of-order packets until the missing packet

- Assume the data from application layer to transport is error free
- Error control unlike flow control involves only the sending and receiving transport layer

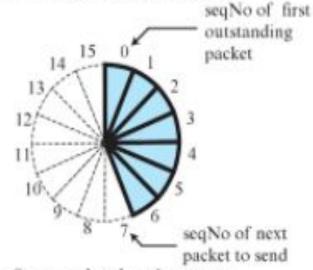


Combination of Error Control and Flow Control

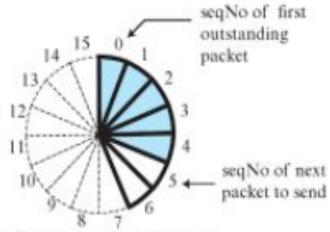
SLIDING Window



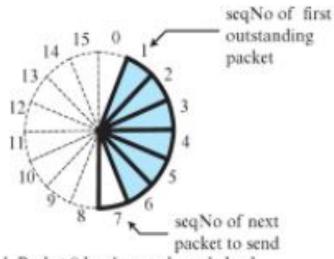
a. Four packets have been sent.



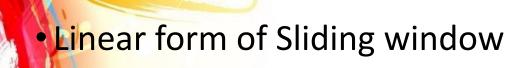
Seven packets have been sent;
 window is full.



b. Five packets have been sent.



 d. Packet 0 has been acknowledged; window slides.





a. Four packets have been sent.



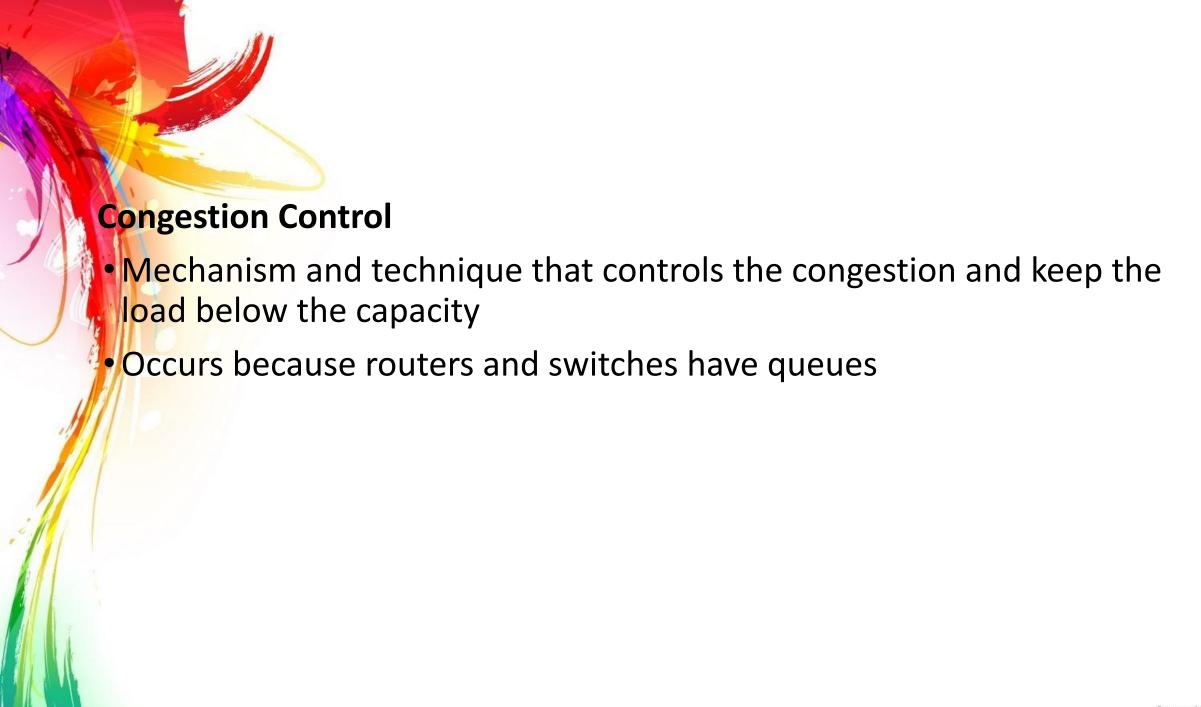
b. Five packets have been sent.



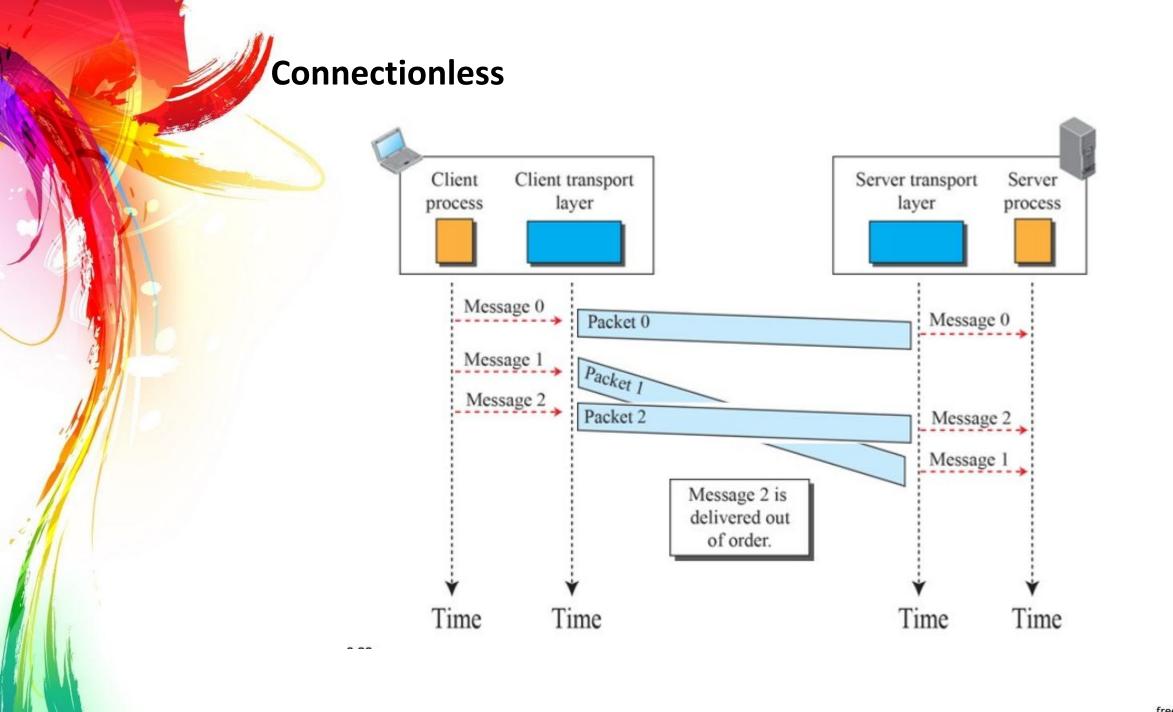
c. Seven packets have been sent;
 window is full.

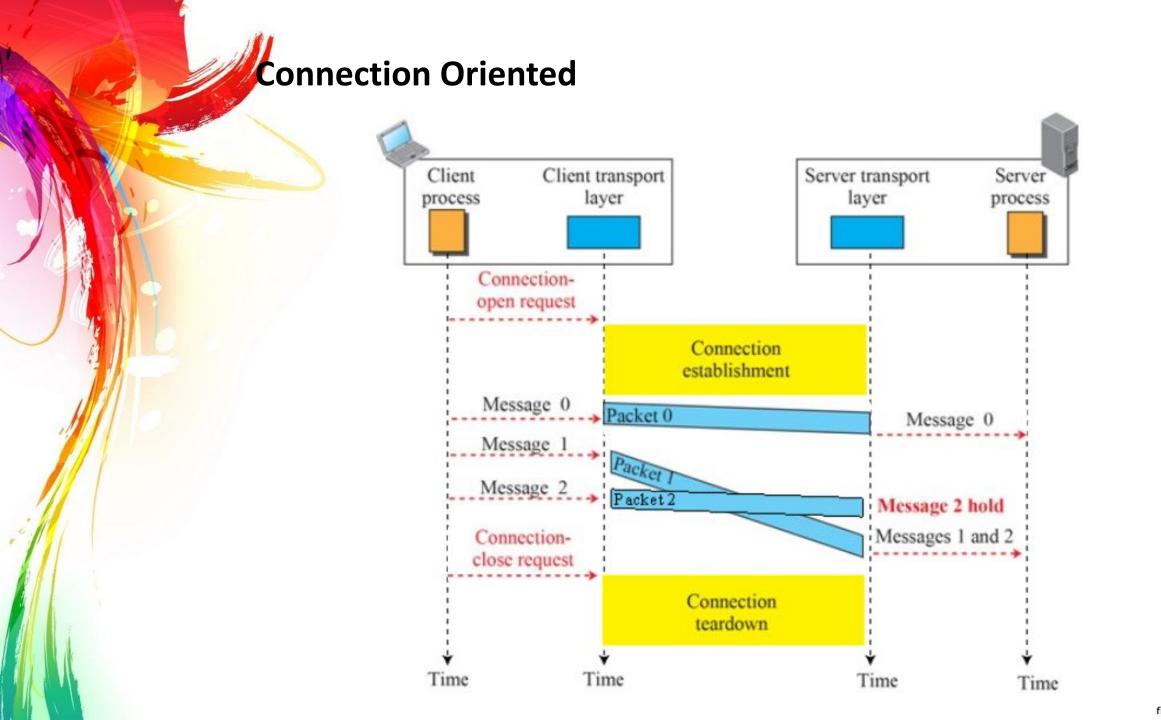


 d. Packet 0 has been acknowledged; window slides.







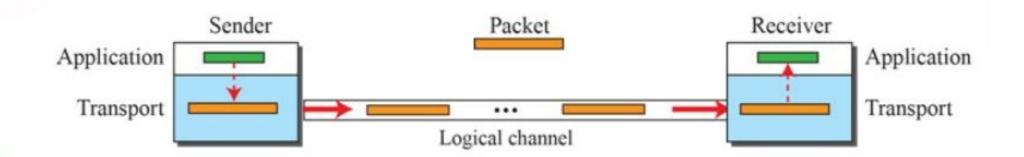


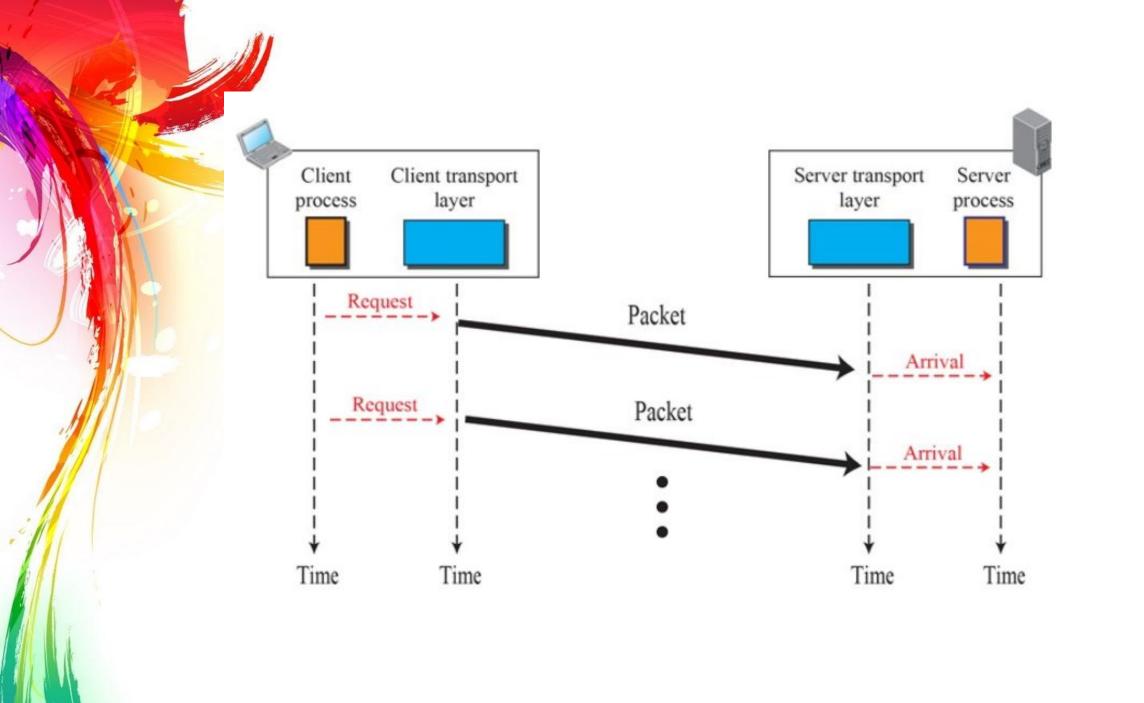
TRANSPORT-LAYER PROTOCOLS

- Simple connectionless protocol
 - No flow control and error control
- Connection-oriented protocol
 - Stop-and-Wait protocol
 - Provide flow and error control
 - Go-Back-N protocol
 - Efficient version of Stop-and-Wait protocol
 - Selective-Repeat Protocol
 - Suited to handle packet loss
 - Piggybacking



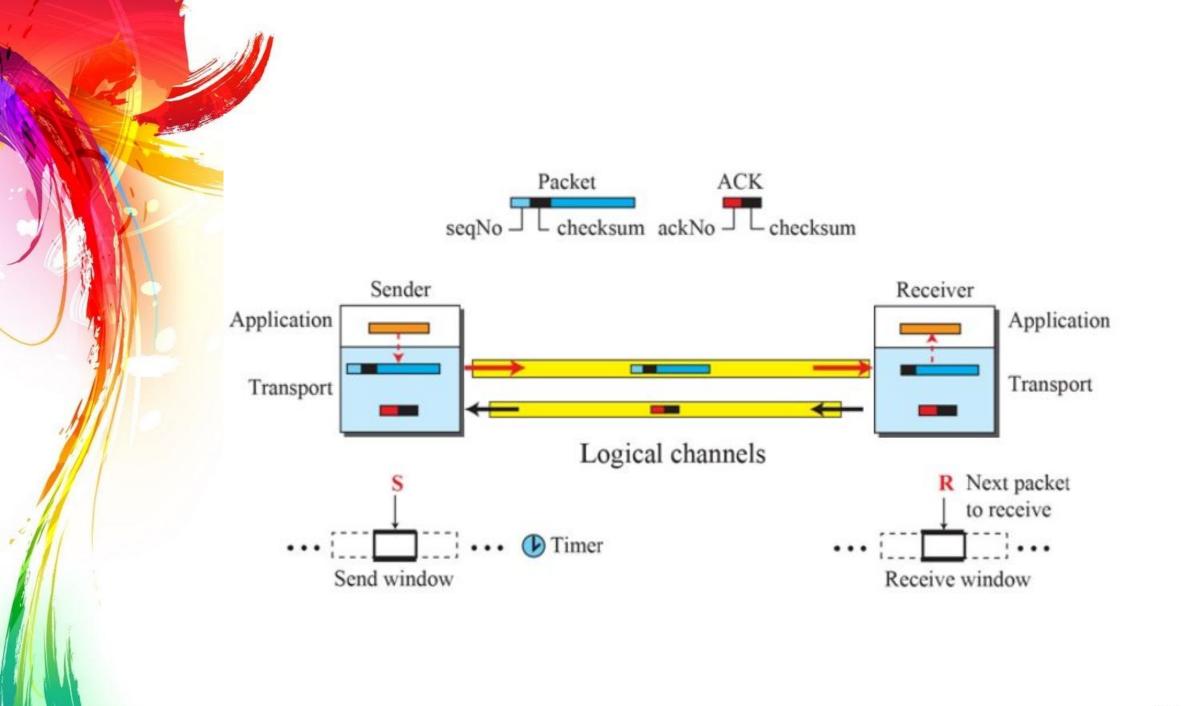
- simple connectionless protocol
- neither flow nor error control
- We assume that the receiver can immediately handle any packet it receives.
- the receiver can never be overwhelmed with incoming packets.





2) Stop-and-Wait Protocol

- connection-oriented protocol
- uses both flow and error control.
- Both the sender and the receiver use a sliding window size 1.
- The sender sends one packet at a time and waits for an acknowledgment before sending the next one.
- To detect corrupted packets, we need to add a checksum to each data packet. When a packet arrives at the receiver site, it is checked. If its checksum is incorrect, the packet is corrupted and silently discarded.

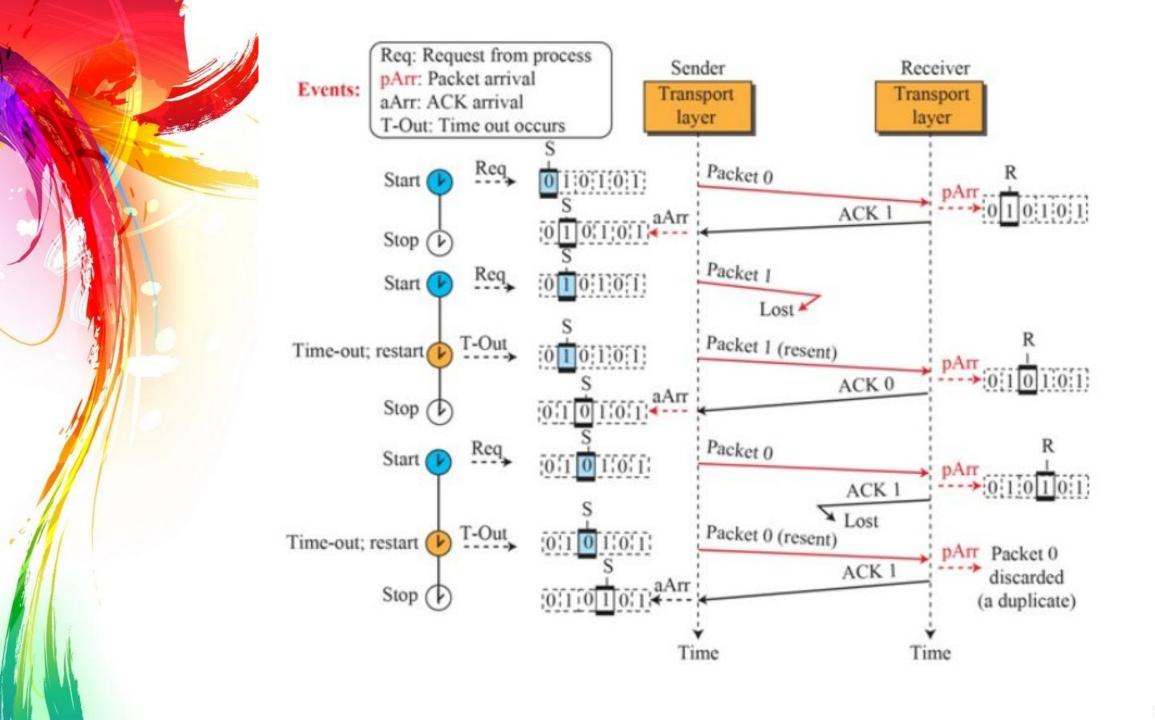


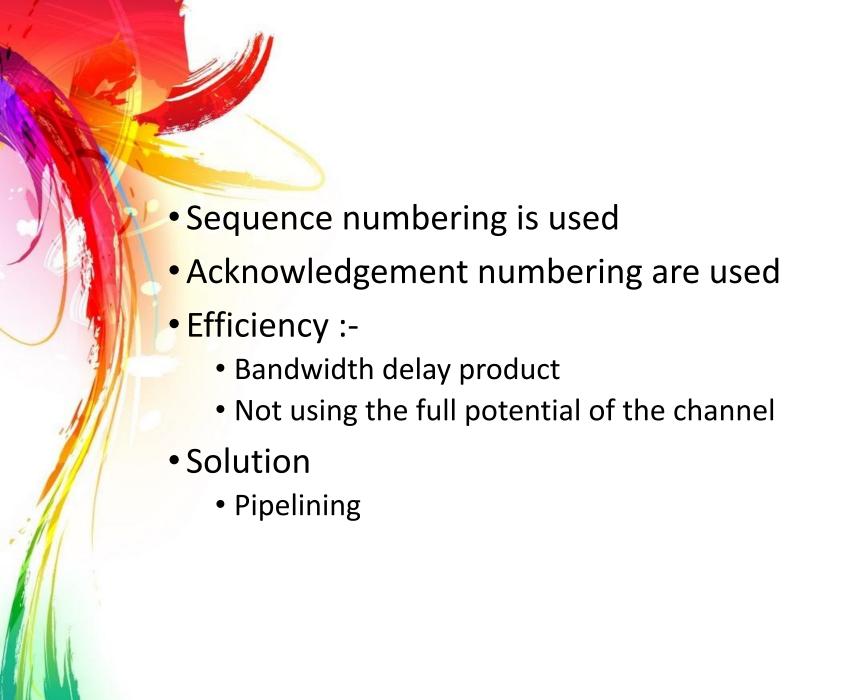
- Example of the Stop-and-Wait protocol.
- Cases:

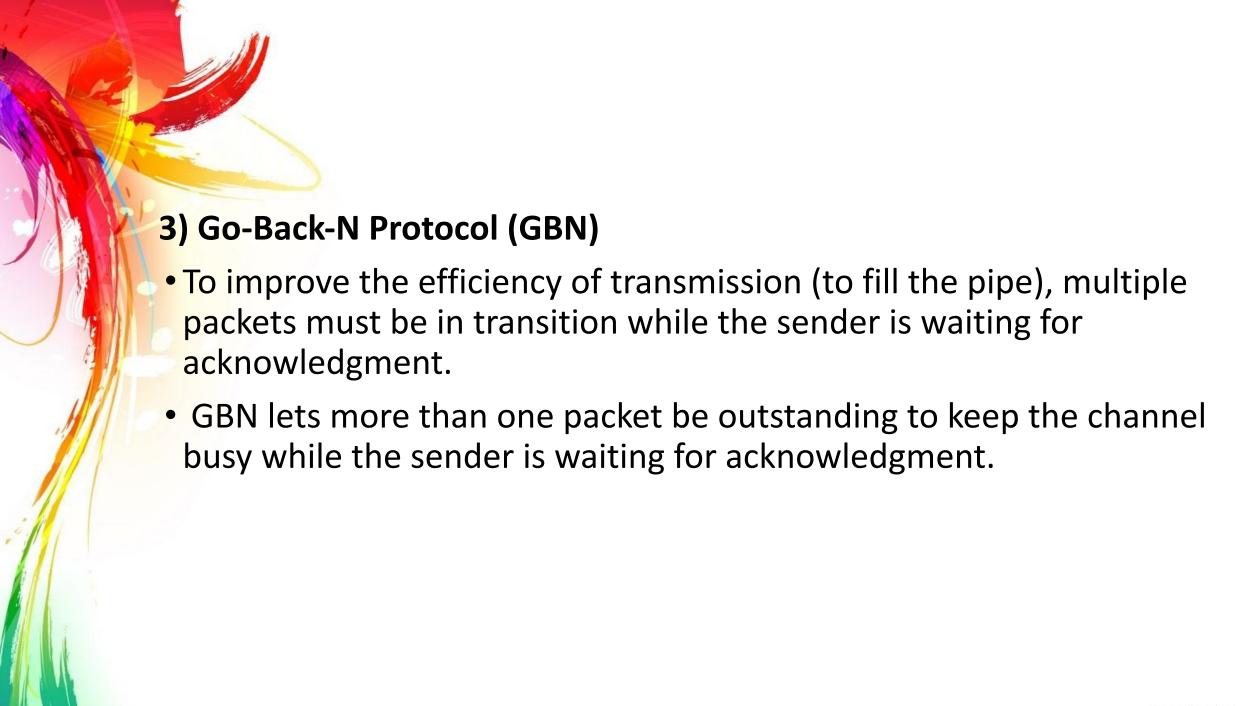
Packet 0 is sent and acknowledged.

Packet 1 is lost and resent after the time-out. The resent packet 1 is acknowledged and the timer stops.

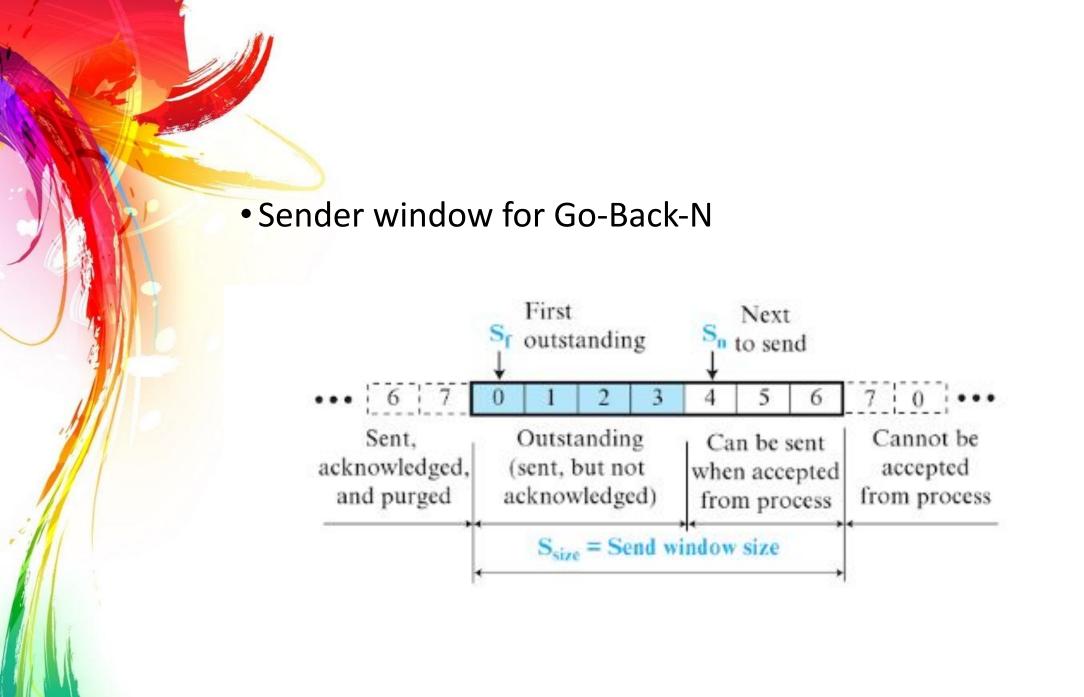
Packet 0 is sent and acknowledged, but the acknowledgment is lost. The sender has no idea if the packet or the acknowledgment is lost, so after the time-out, it resends packet 0, which is acknowledged.

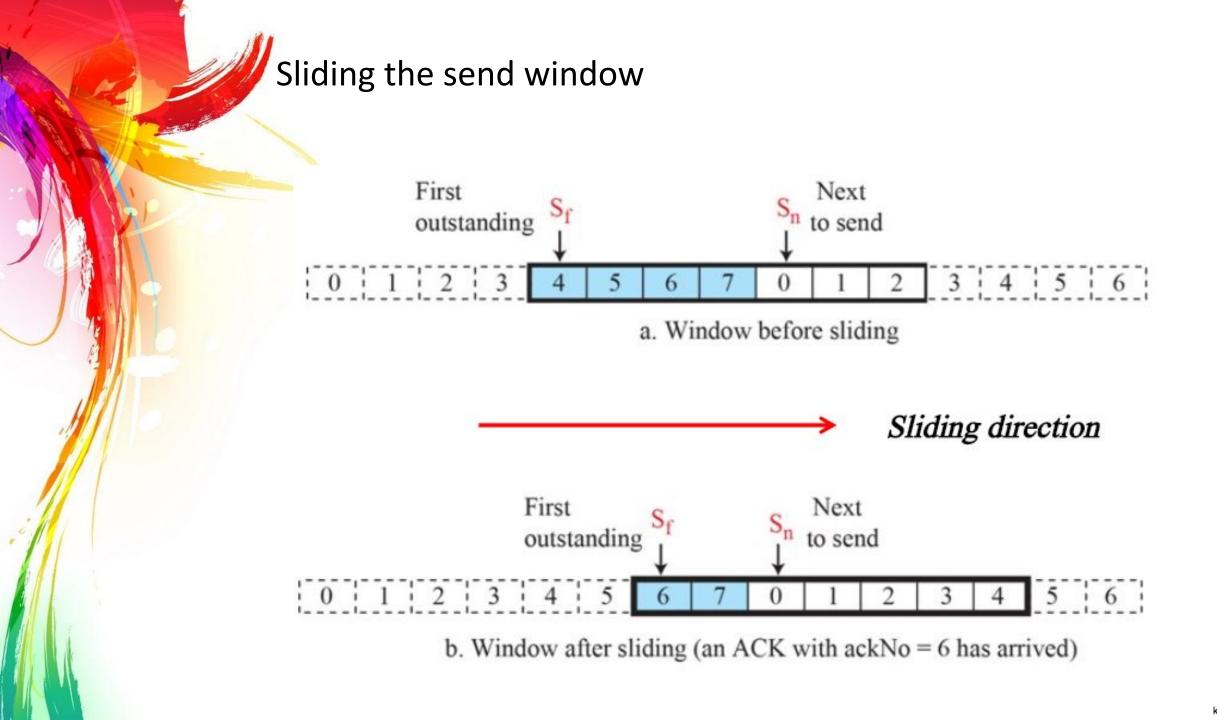






Send window for Go-Back-N Packet ACK checksum ackNo - checksum Sender Receiver Application Application Transport Transport Logical channels R Next packet S_fFirst outstanding S_n Next 1 to send (I) Timer to receive Receive window Send window

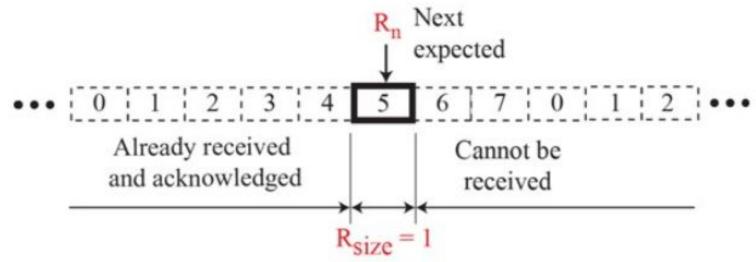


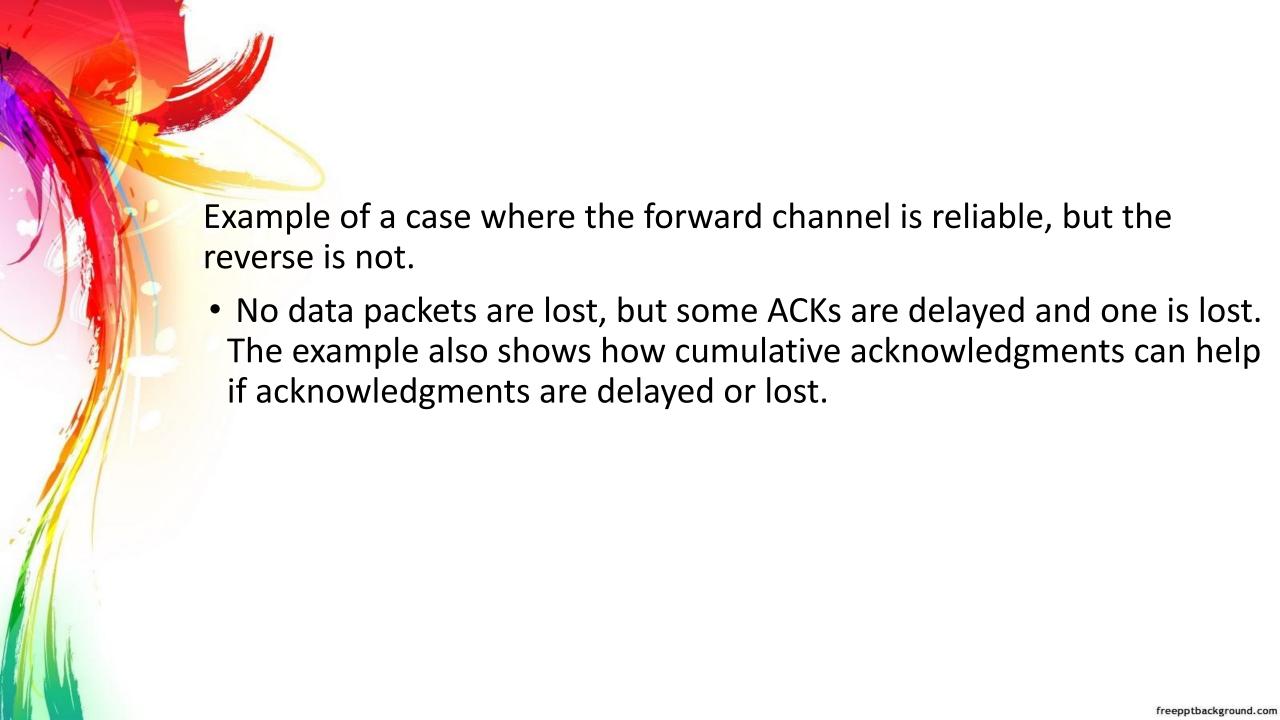


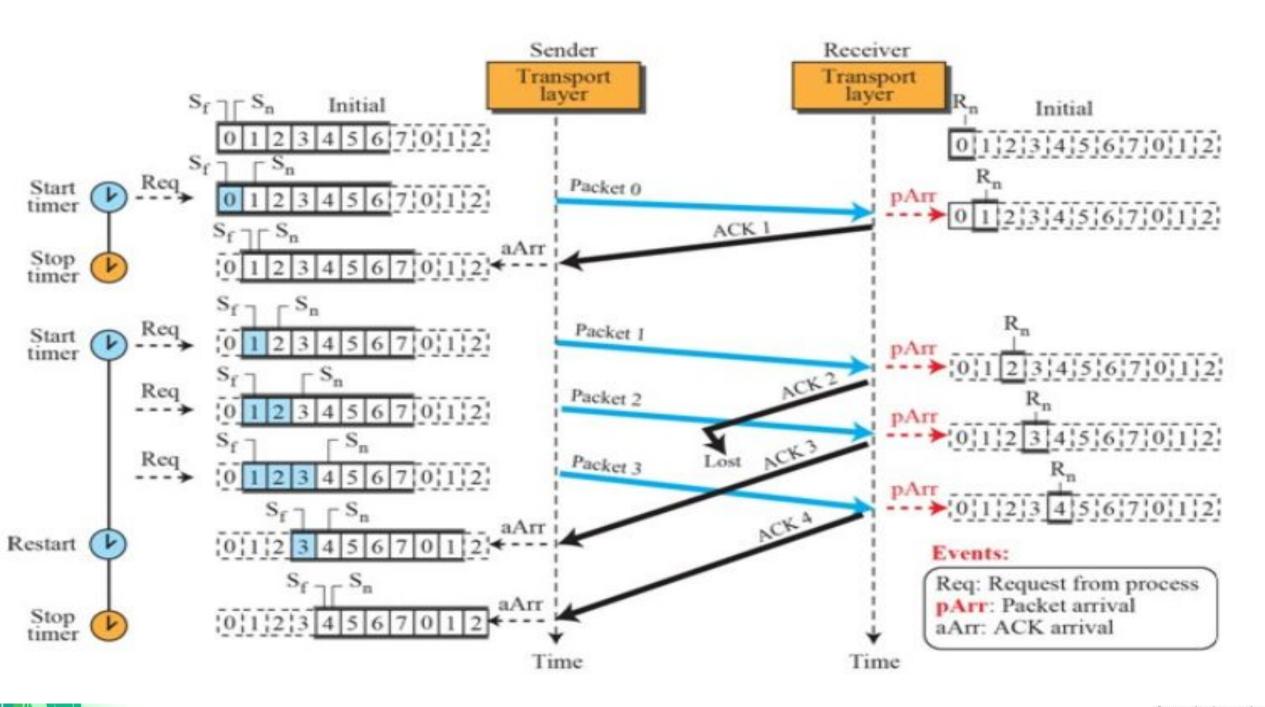


Receive window for Go-Back-N

- Size of the receiving window is always 1
- Any packet coming out of order is discarded and resent (timer)

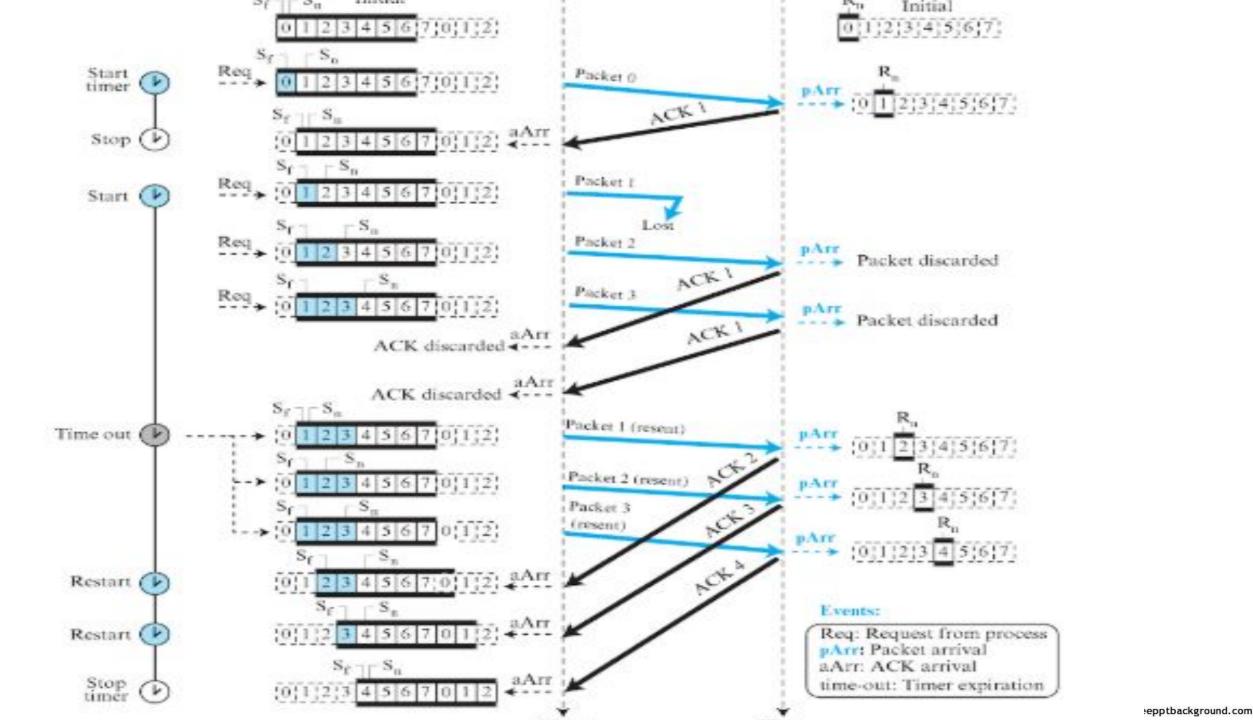


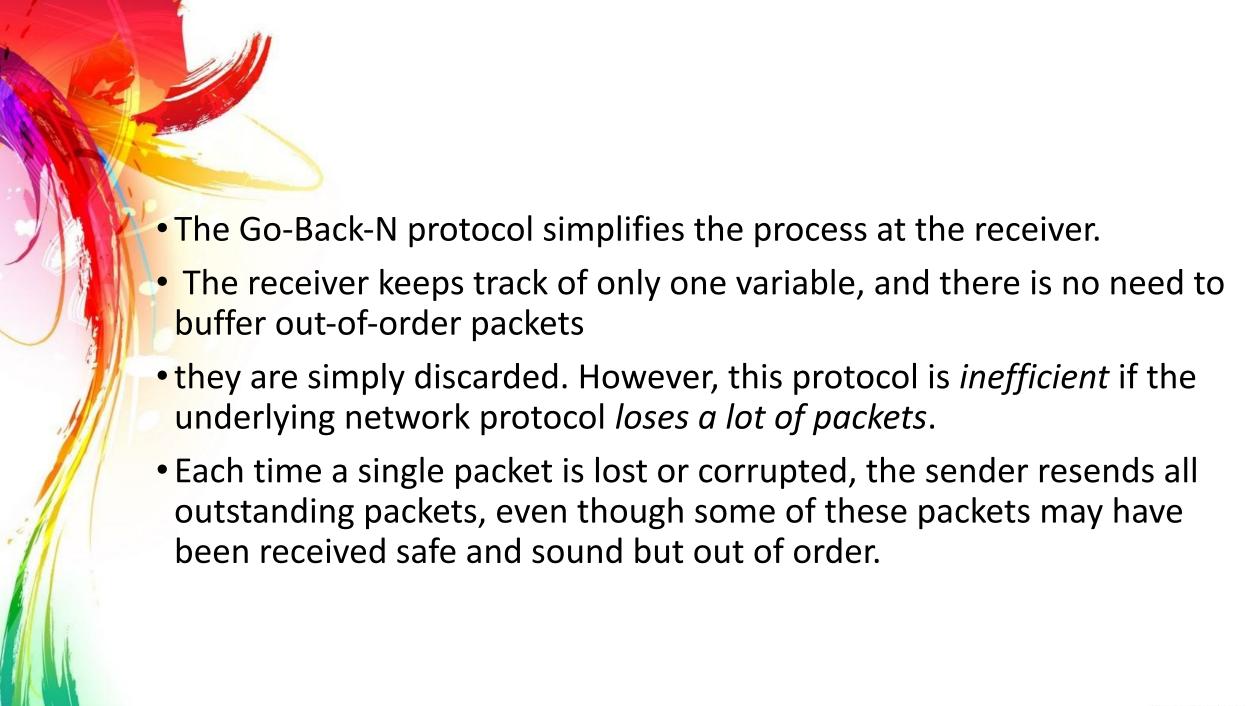


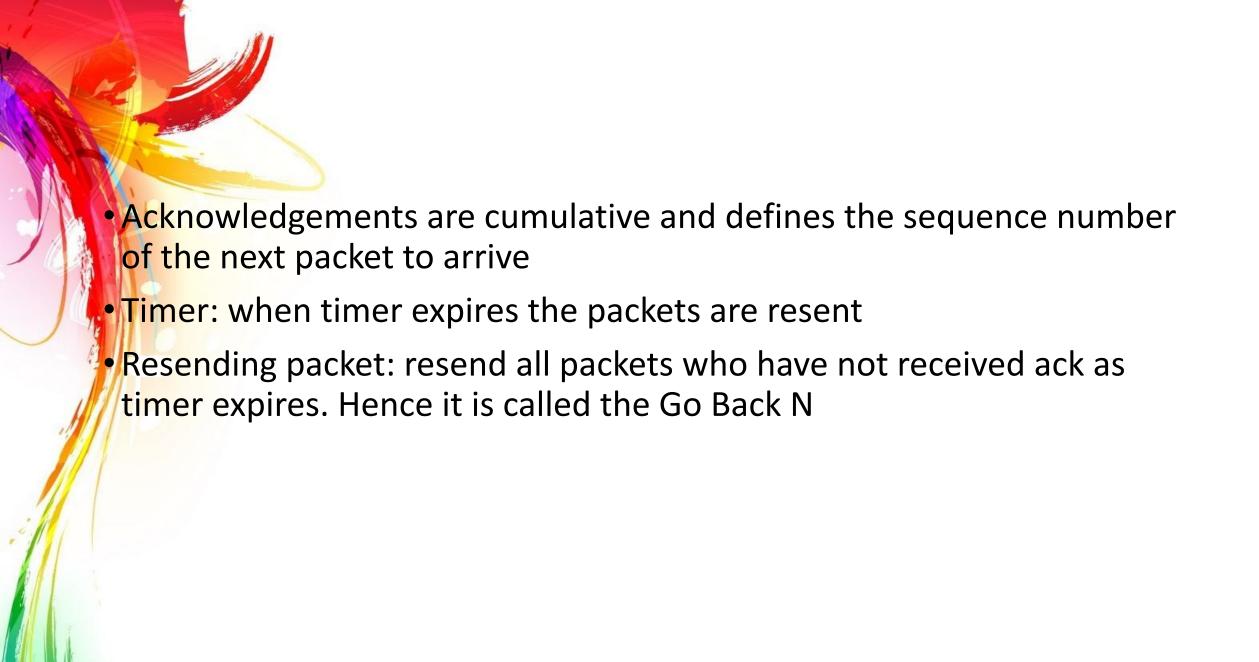


freepptbackground.com

Next shows what happens when a packet is lost. Packets 0, 1, 2, and 3 are sent. However, packet 1 is lost. The receiver receives packets 2 and 3, but they are discarded because they are received out of order (packet 1 is expected). When the receiver receives packets 2 and 3, it sends ACK1 to show that it expects to receive packet 23. However, these ACKs are not useful for the sender because the ackNo is equal to S f, not greater that S f. So the sender discards them. When the time-out occurs, the sender resends packets 1, 2, and 3, which are acknowledged.

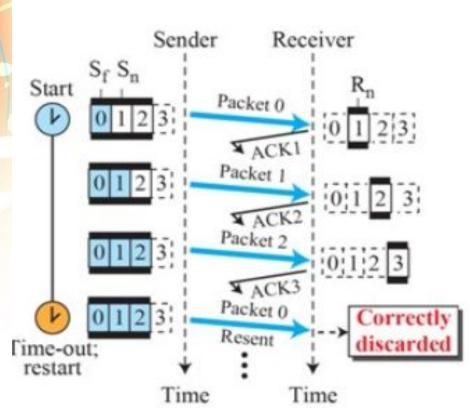


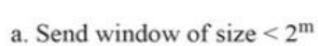


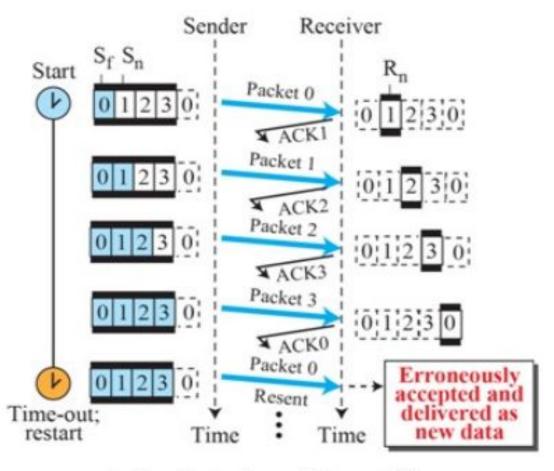


Send window size for Go-Back-N

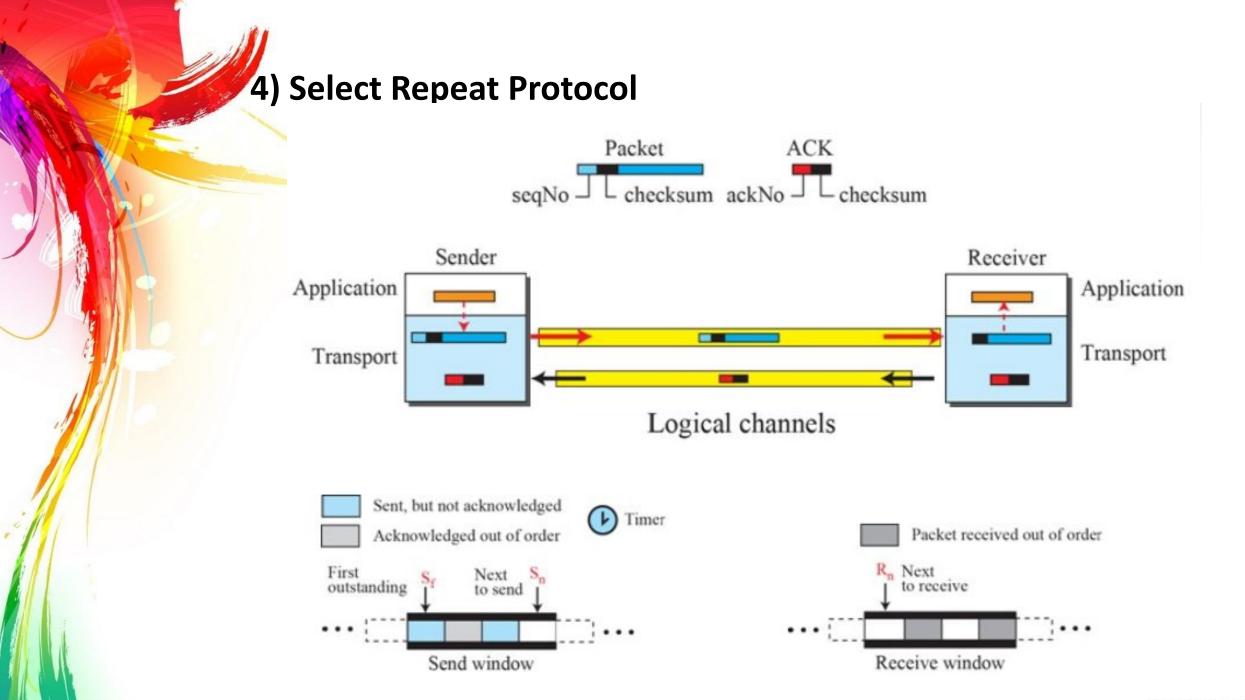
• Size of the send window must be < 2^m



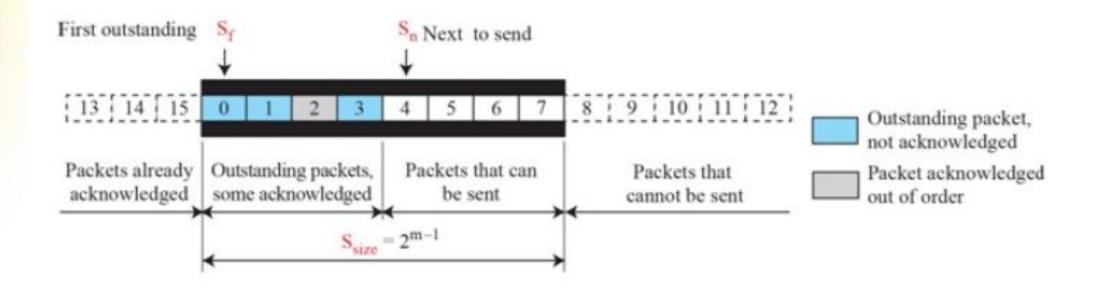


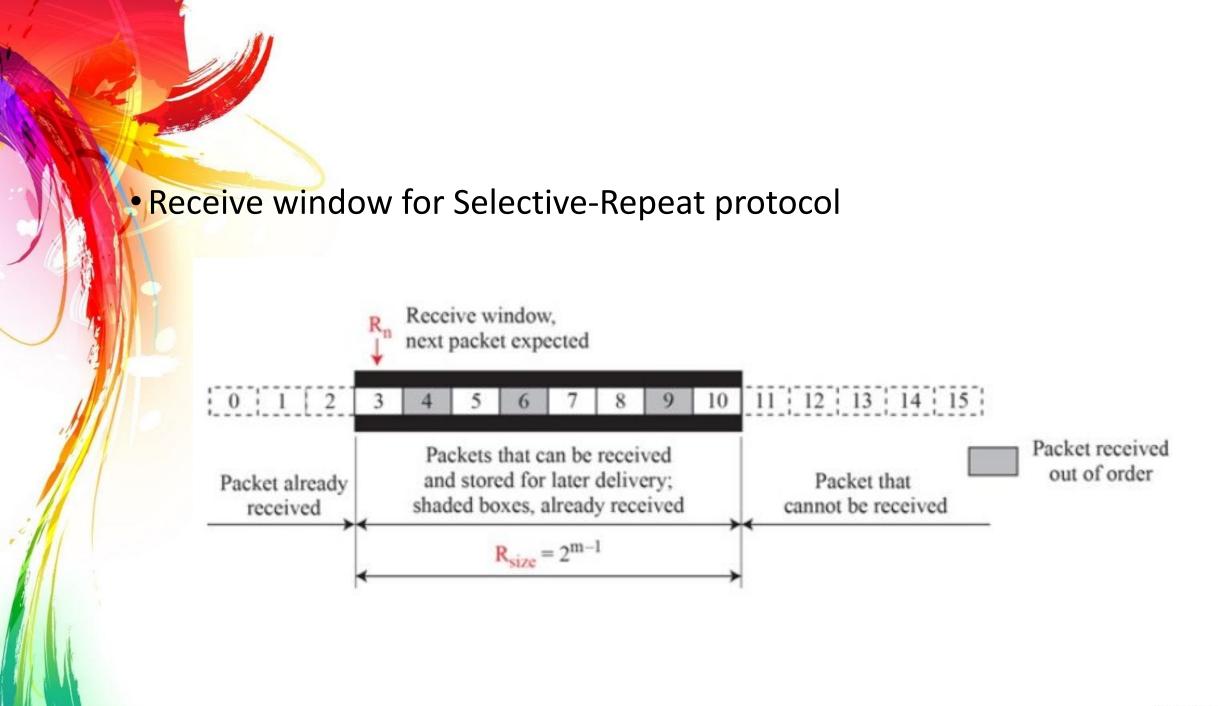


b. Send window of size = 2^{m}

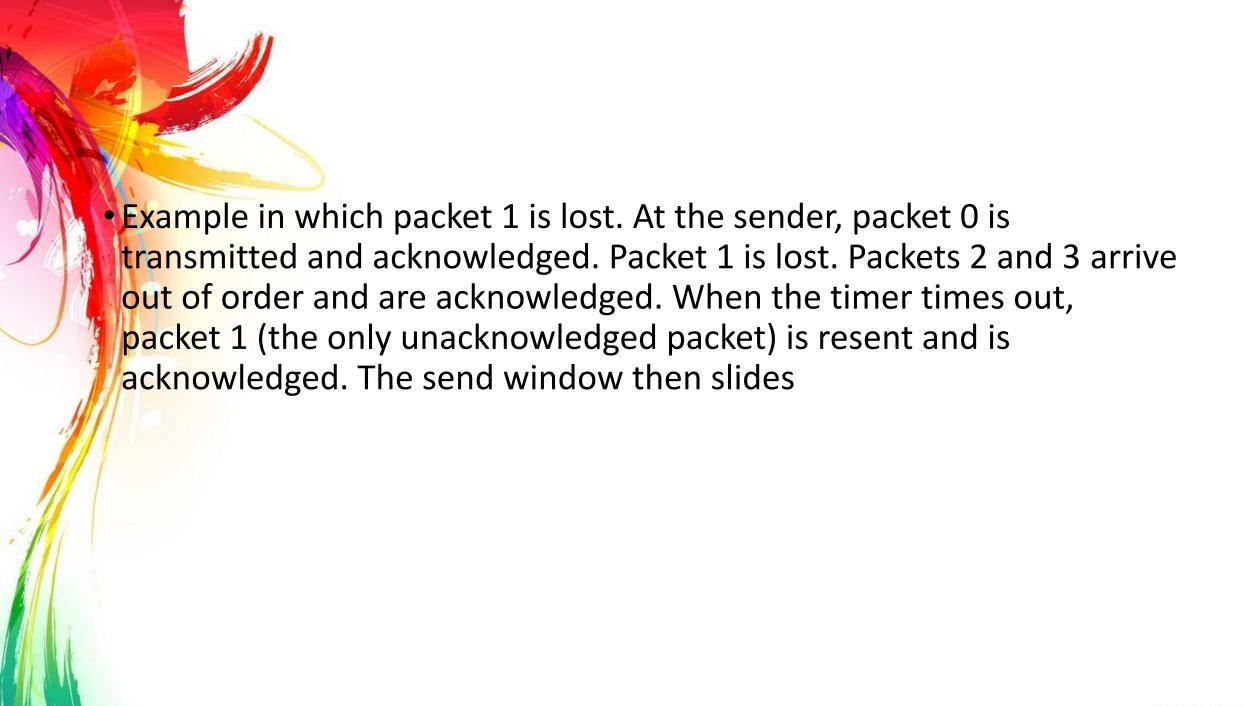


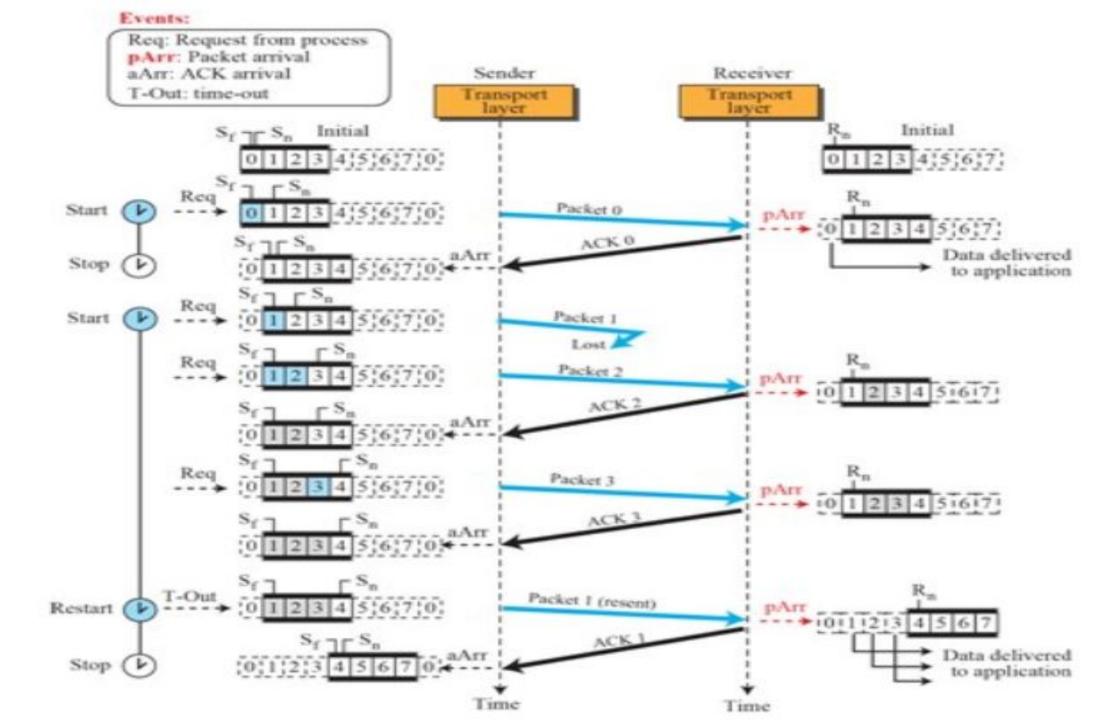
Send window for Selective-Repeat protocol





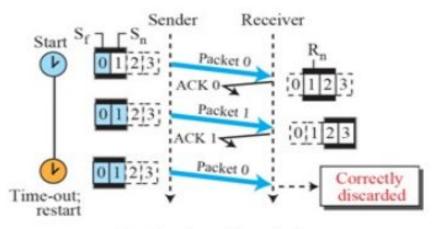




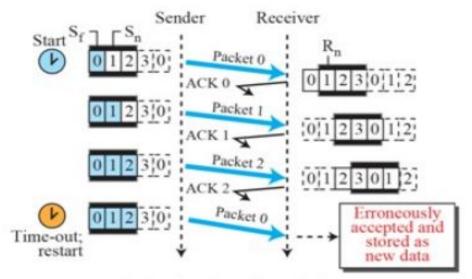


- At the receiver site we need to distinguish between the acceptance of a packet and its delivery to the application layer.
- At the second arrival, packet 2 arrives and is stored and marked (shaded slot), but it cannot be delivered because packet 1 is missing.
- At the next arrival, packet 3 arrives and is marked and stored, but still none of the packets can be delivered. Only at the last arrival, when finally a copy of packet 1 arrives, can packets 1, 2, and 3 be delivered to the application layer.
- There are *two conditions* for the delivery of packets to the application layer:
 - First, a set of consecutive packets must have arrived.
 - Second, the set starts from the beginning of the window.

Selective-Repeat, window size



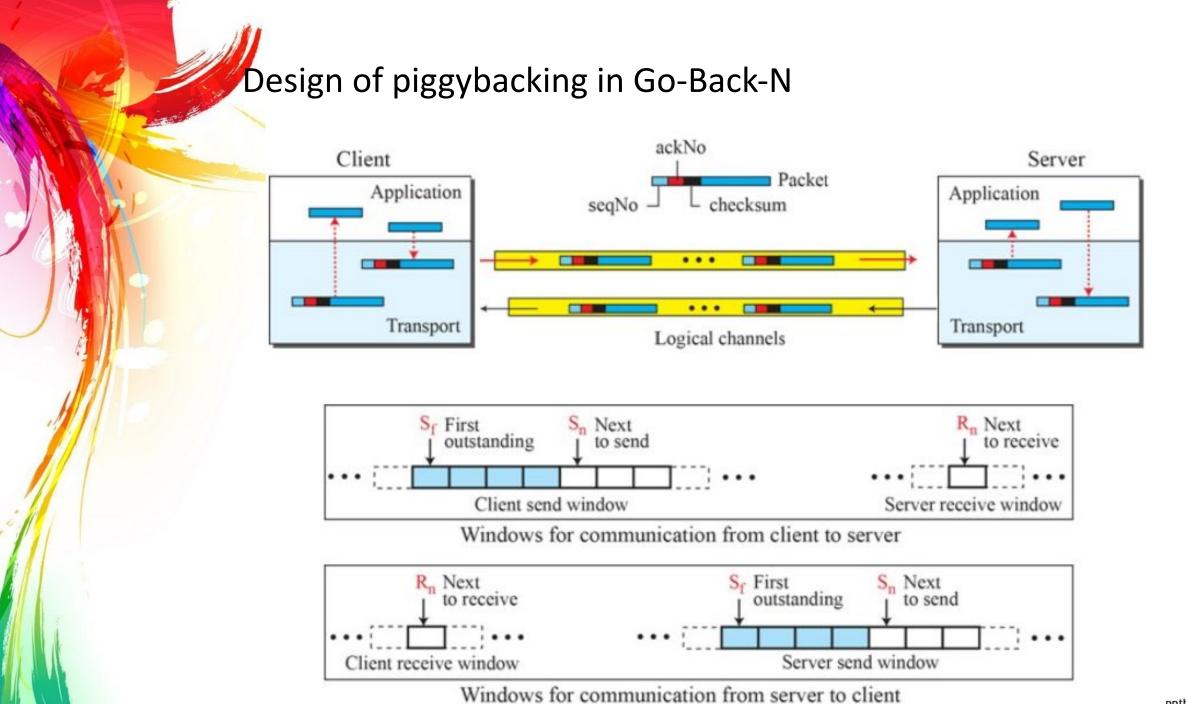
a. Send and receive windows of size = 2^{m-1}



b. Send and receive windows of size > 2^{m-1}

5) Bidirectional Protocols: Piggybacking

- All discussed so far are all unidirectional: data packets flow in only one direction and acknowledgments travel in the other direction.
- In real life, data packets are normally flowing in both directions: from client to server and from server to client.
- This means that acknowledgments also need to flow in both directions.
- A technique called *piggybacking* is used to improve the efficiency of the bidirectional protocols.
- When a packet is carrying data from A to B, it can also carry acknowledgment feedback about arrived packets from B



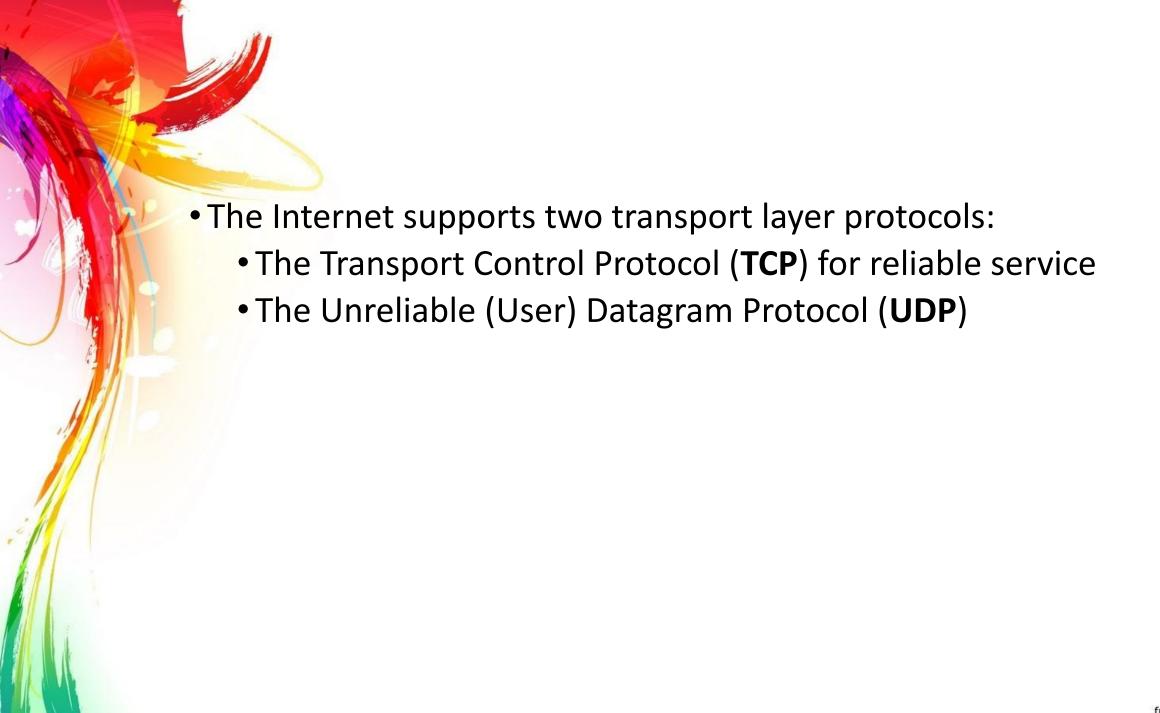
Transport layer duties/ Design Criterion

- Addressing
 - Port numbers to identify which network application
- Reliability
 - Flow control
 - Error Control
- Connection control
 - Connection-oriented
 - Connectionless
- Packetizing

process

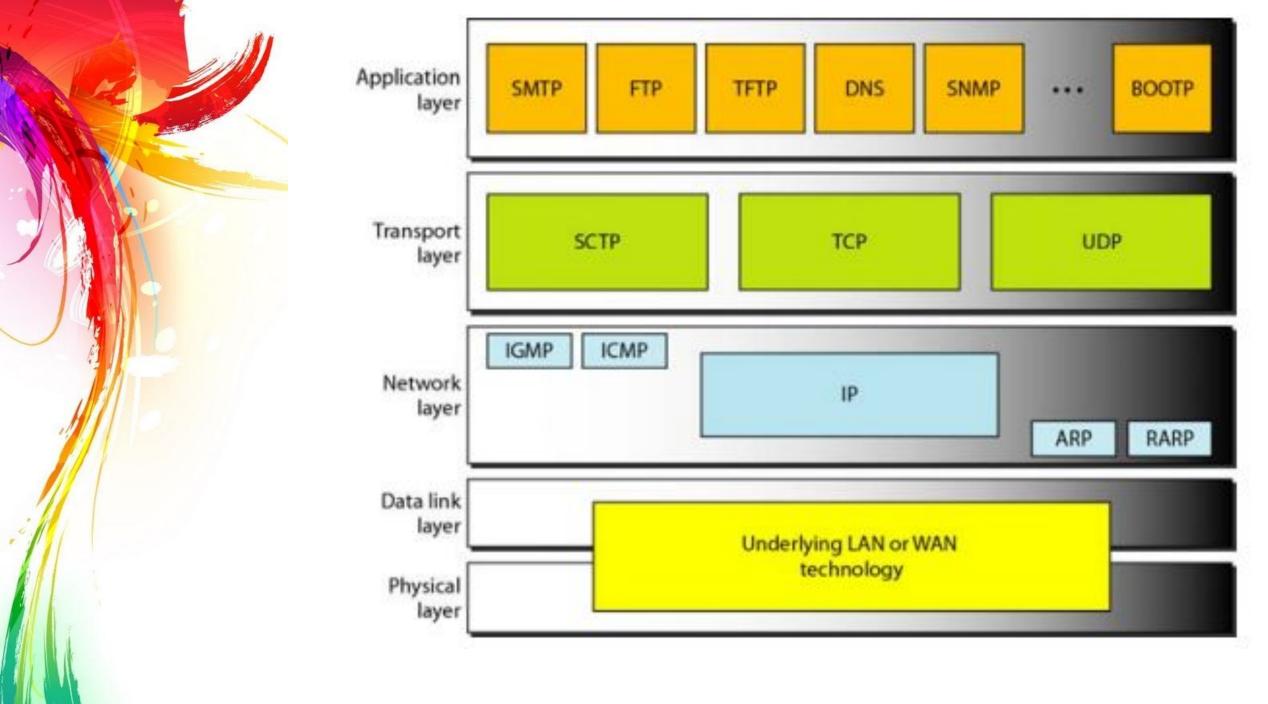
- Sender side: breaks **application messages** into segments, passes them to network layer
- Transport layer at the receiving host deliver data to the receiving

epptbackground.com



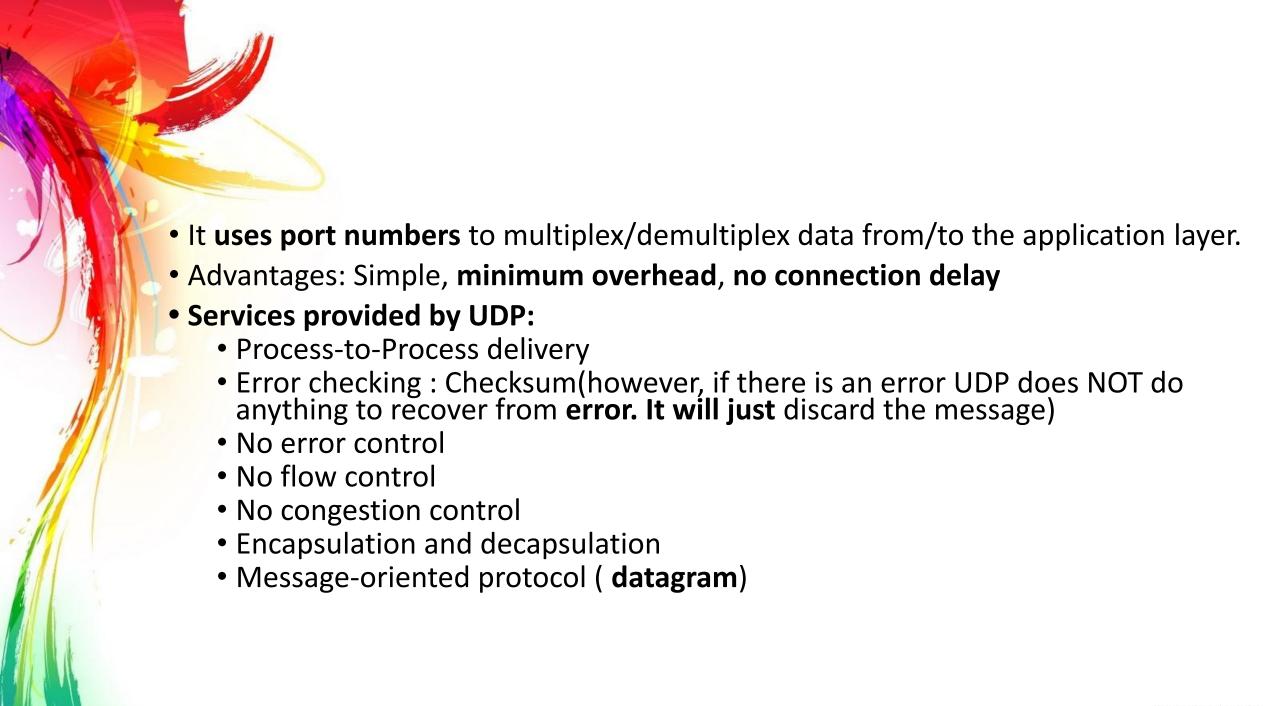
Processes communicating across network

- Process is an instance of a program in execution.
- Processes on two hosts communicate with each other by sending and receiving messages
- The process receives messages from, and sends messages into the network through its socket
- A <u>socket</u> is the **interface** between the **application layer** and the **transport layer** within a host.
- **Sockets** are the **programming interface** used to <u>build network</u> <u>applications</u> over the internet.
- Programmers can select <u>which transport layer protocol</u> (UDP or TCP) to be used by the application and select few transport-layer parameters (maximum buffer size, Maximum segment size, starting sequence number of segment).



User Datagram Protocol (UDP)

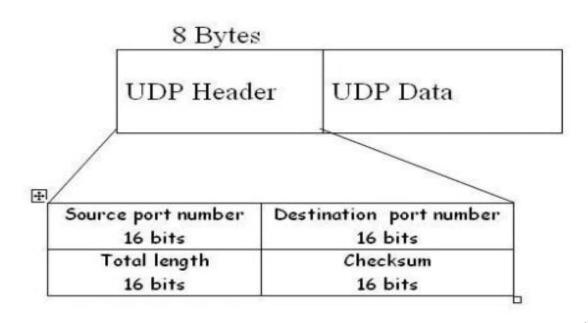
- Connectionless
 - No handshaking between UDP sender, receiver
 - Each UDP segment handled independently of others
- Unreliable protocol has no flow and error control
 - A UDP segment can be lost, arrive out of order, duplicated, or corrupted
 - Checksum field checks error in the entire UDP segment. It is Optional
 - UDP doe not do anything to recover from an error it simply discard the segment
 - Application accepts full responsibility for errors

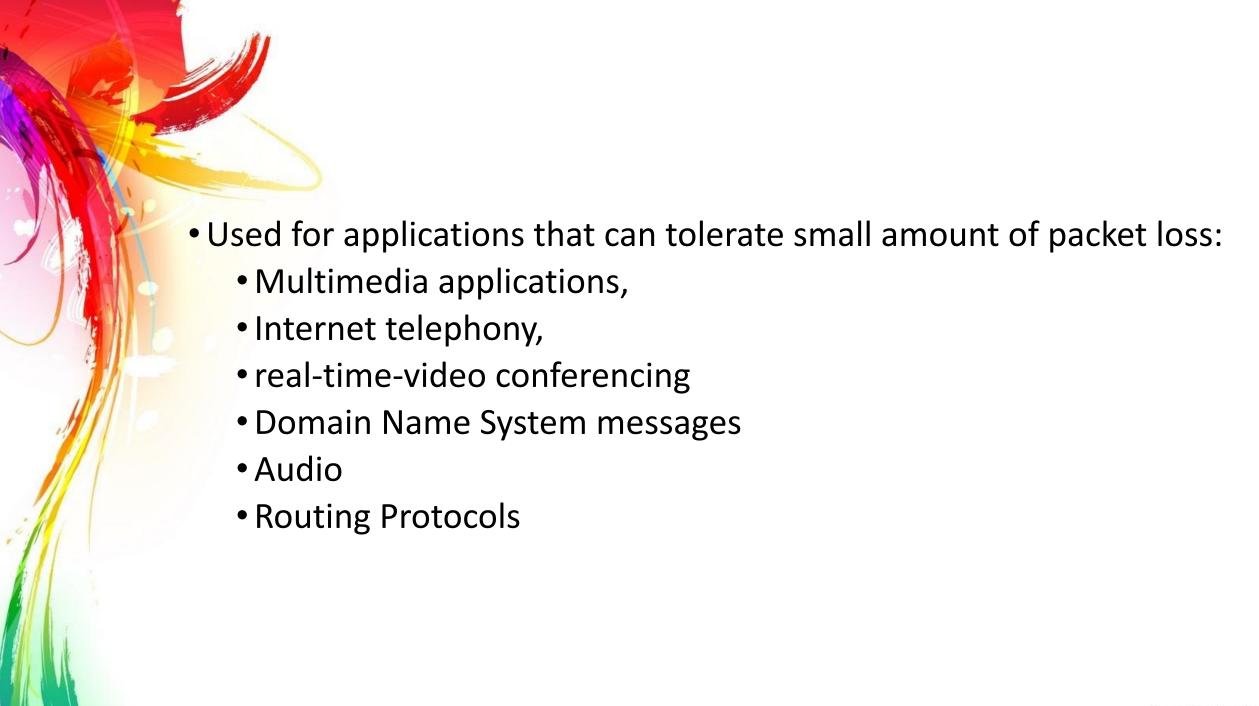


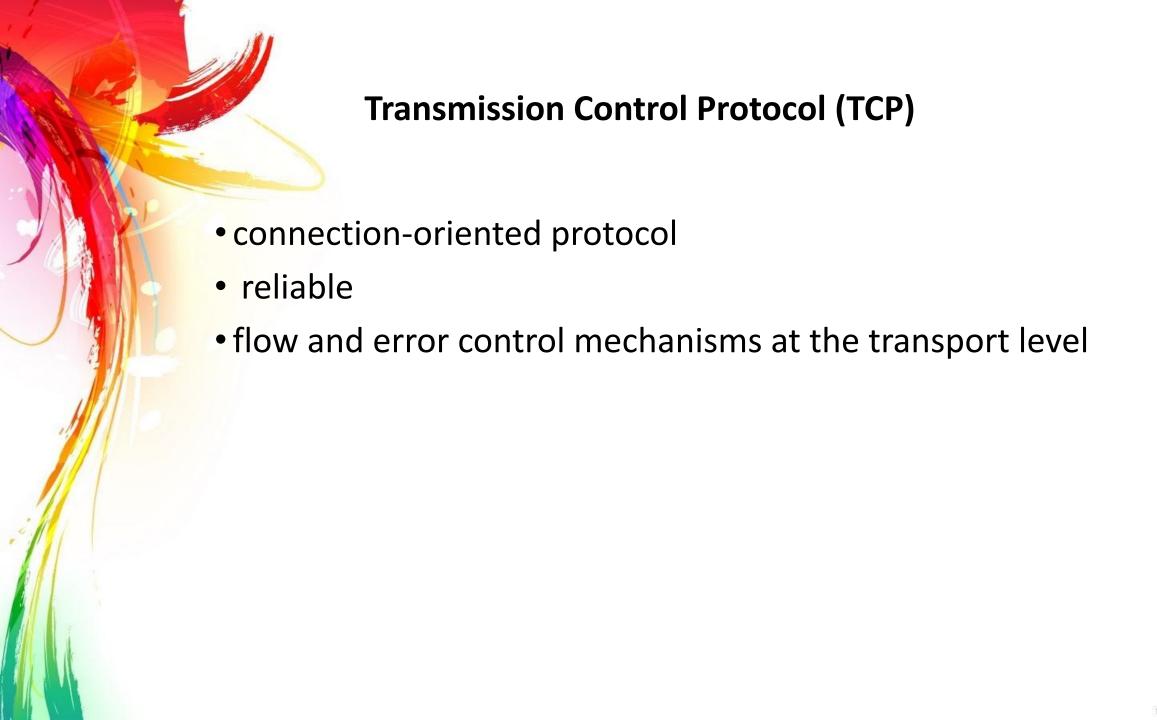
User datagram format

- Header size = 8 bytes
 Minimum UDP process data size 0 bytes
- Maximum UDP process data size=

65535 – 20 (network layer headers) - 8 (UDP headers) = 65507 bytes





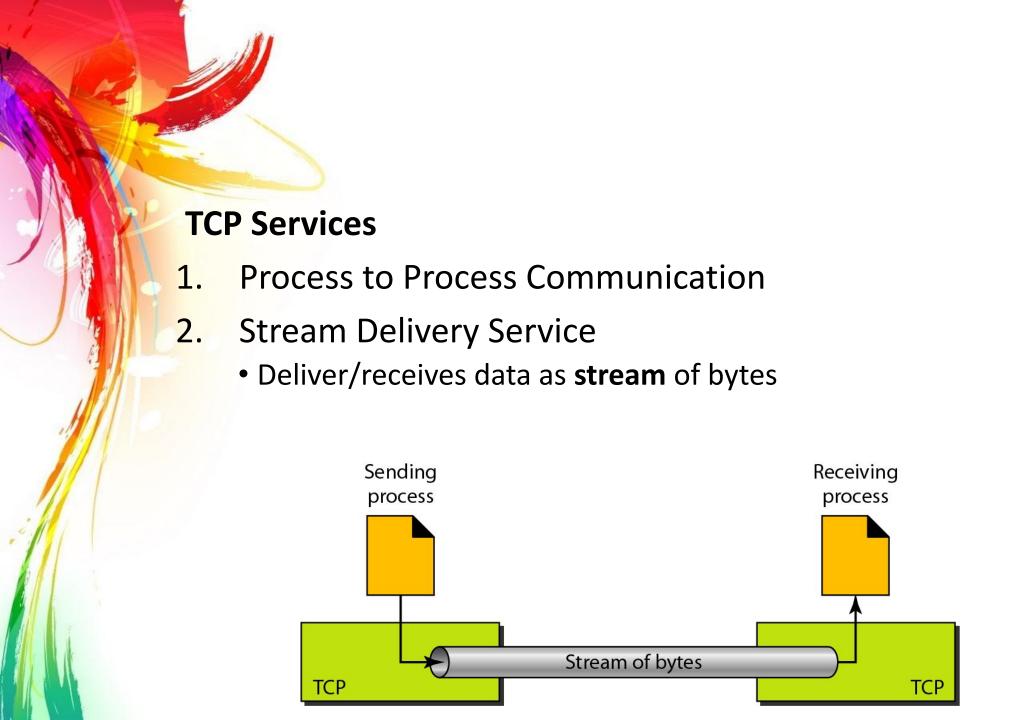


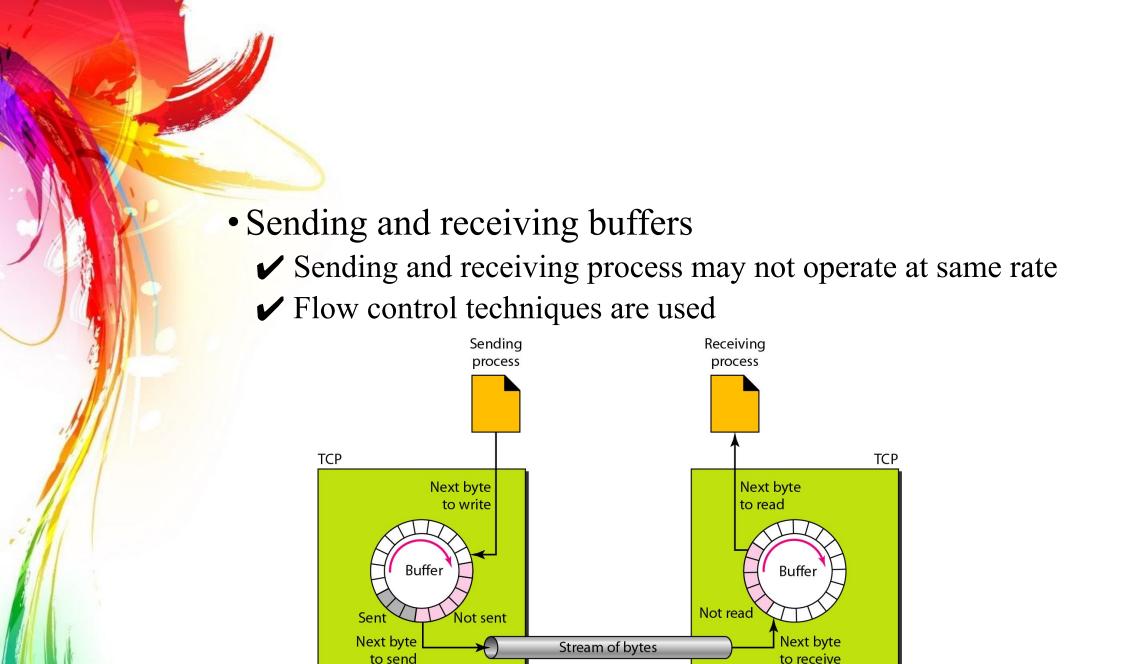


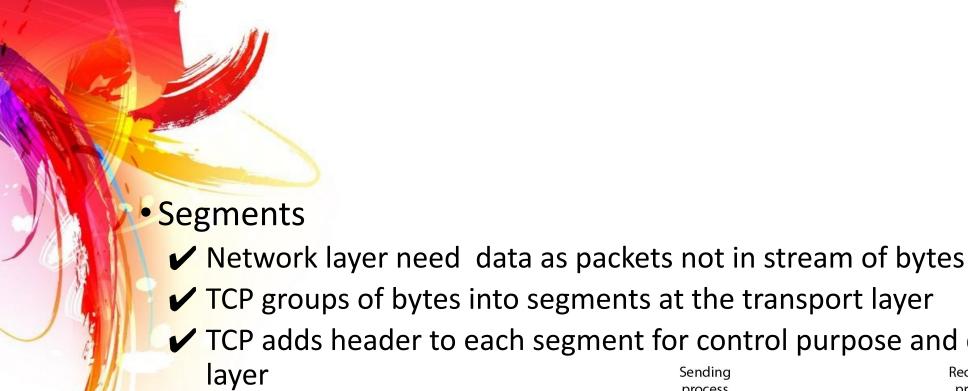
Well-known ports used by TCP

Port	Protocol	Description
7	Echo	Echoes a received datagram back to the sender
9	Discard	Discards any datagram that is received
11	Users	Active users
13	Daytime	Returns the date and the time
17	Quote	Returns a quote of the day
19	Chargen	Returns a string of characters
20	FTP, Data	File Transfer Protocol (data connection)
21	FTP, Control	File Transfer Protocol (control connection)
23	TELNET	Terminal Network
25	SMTP	Simple Mail Transfer Protocol
53	DNS	Domain Name Server
67	BOOTP	Bootstrap Protocol
79	Finger	Finger
80	HTTP	Hypertext Transfer Protocol
111	RPC	Remote Procedure Call

23.7



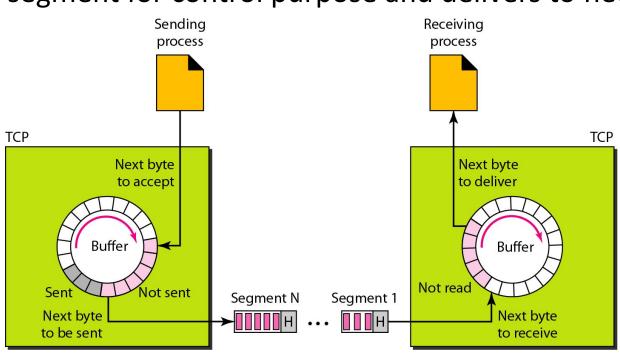




✓ TCP groups of bytes into segments at the transport layer

✓ TCP adds header to each segment for control purpose and delivers to network

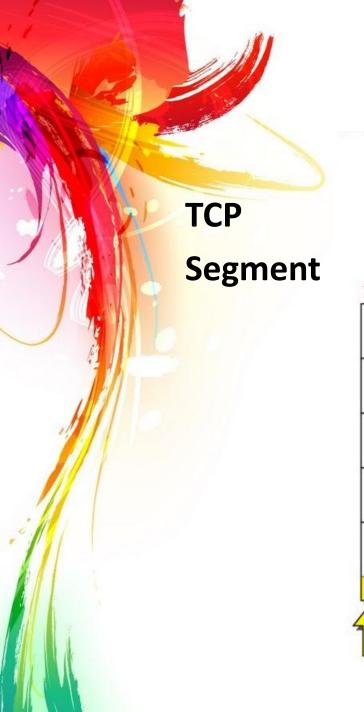
✓ Segment size can be variable

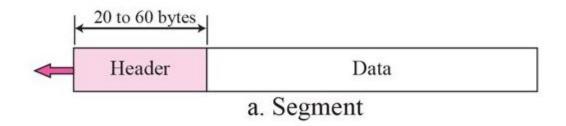


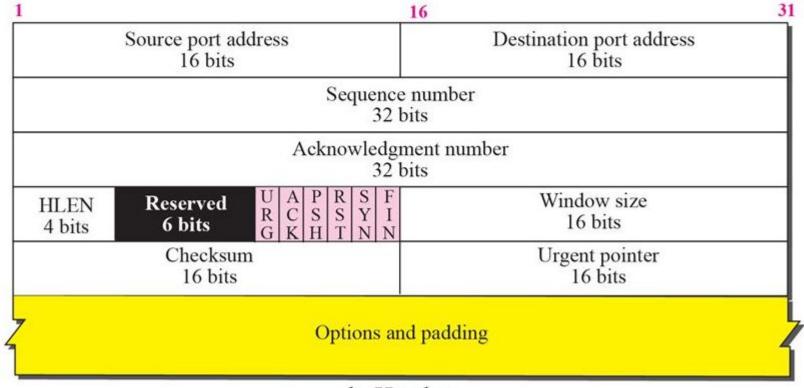


TCP Numbering System

- ✓ Byte no not segment no
 - Very octet is numbered
 - Numbering is independent at sender and receiver
 - Need not start with 0, it starts with arbitrary no
- ✔ For tracking
 - Sequence no
 - To each segment
 - Sequence no of 1st segment is referred to as ISN (first byte)
 - Next segment is previous no +1
 - Acknowledgement no
 - Sequence no +1 *
 - (* one block of segment)









URG: Urgent pointer is valid

ACK: Acknowledgment is valid

PSH: Request for push

RST: Reset the connection

SYN: Synchronize sequence numbers

FIN: Terminate the connection



Description of flags in the control field

Flag	Description	
URG	The value of the urgent pointer field is valid.	
ACK	The value of the acknowledgment field is valid.	
PSH	Push the data.	
RST	Reset the connection.	
SYN	Synchronize sequence numbers during connection.	
FIN	Terminate the connection.	

- TCP segment encapsulates data received from application layer
- itself be encapsulated in an IP datagram at the network layer
- later as frame in data link layer

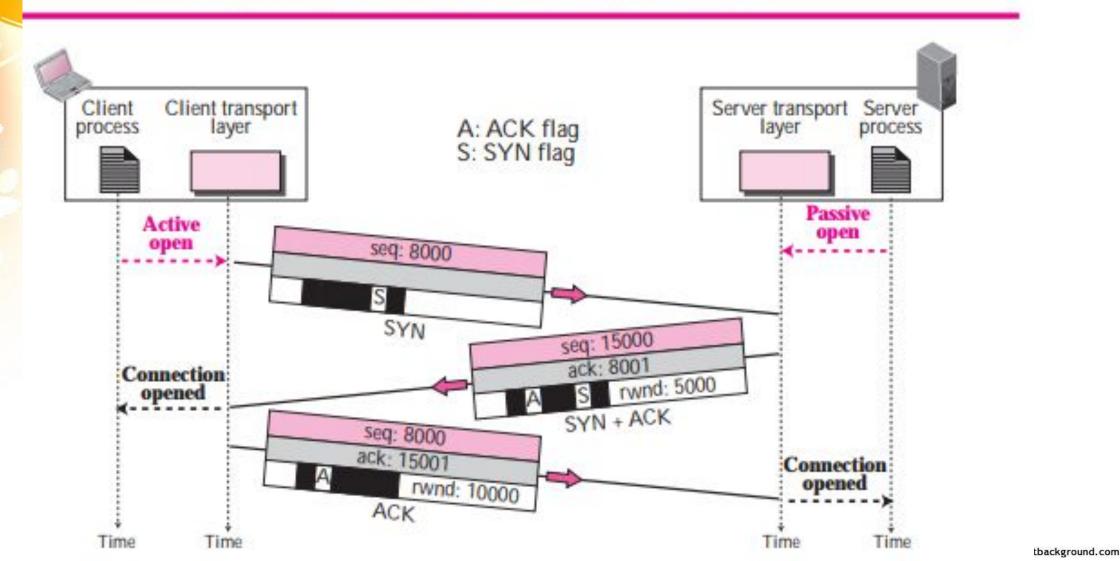
TCP Connection

Connection oriented

- Both sender and receiver need to establish connection separately
- Logical path b/w transport layer facilitates the message and acknowledgement delivery and retransmission of packets.
- 3 phases
 - 1. Establishment of logical connection
 - 2. Data *exchange* in both direction
 - 3. Connection *Termination*
- Fully duplex mode

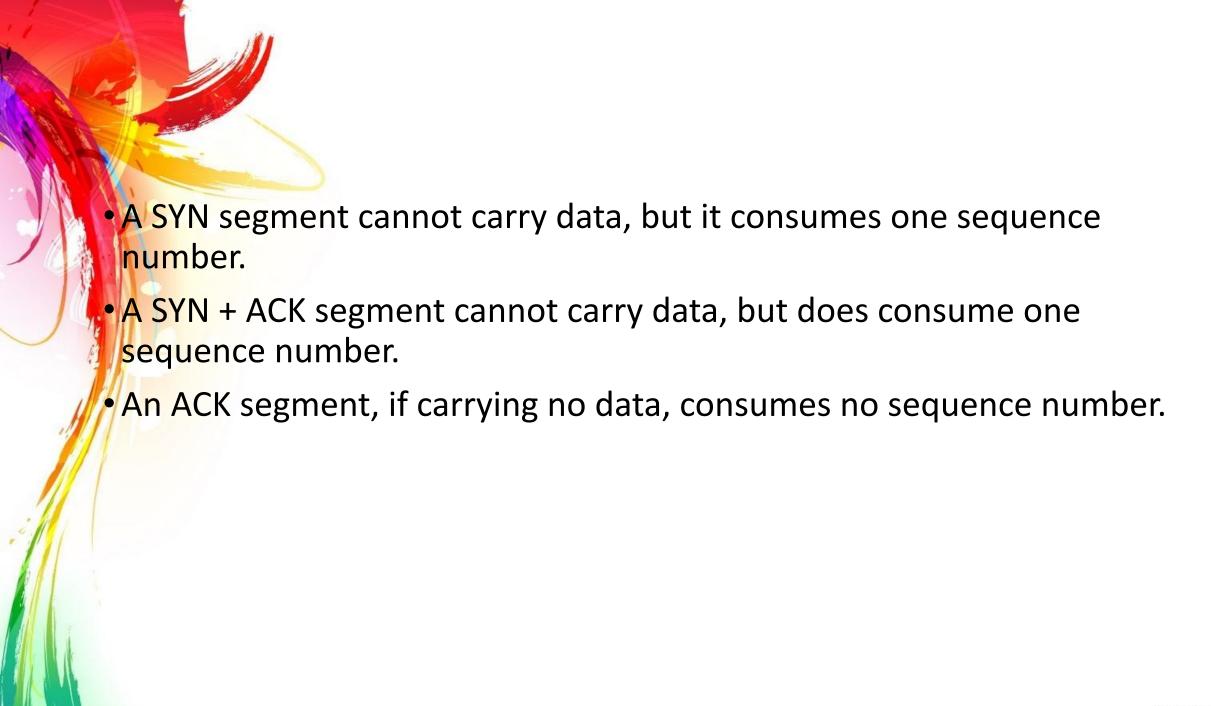
Step 1: Establishment of logical connection

Connection establishment using three-way handshaking

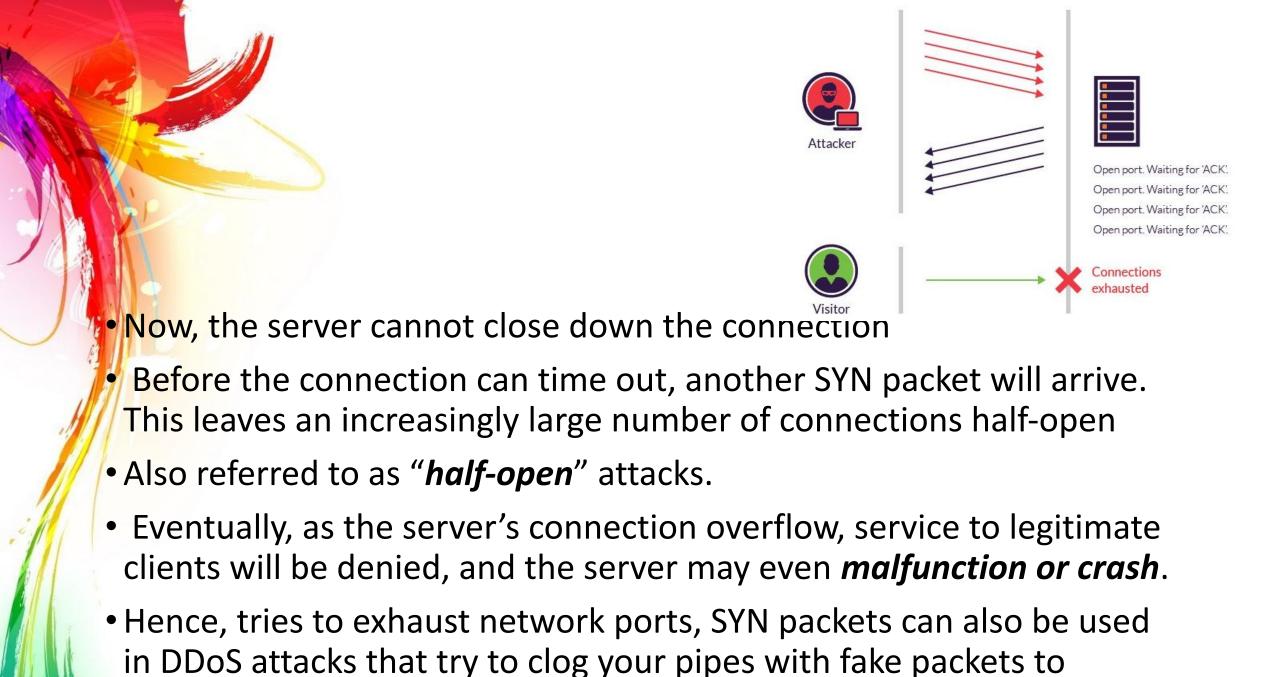


3 way handshake

- Passsive open server ready to accept connections
- Active open client that wishes to connect to open server.
- SYN segement to synchronize sequence nos
 - First no is called ISN (Initial Sequence No)
 - No data
 - No other flags or window size
- SYN + ACK
 - Dual purpose
 - Server sends ISN and ACK to the SYN received
 - As ACK, it need to send the receive window size (*rwnd*)



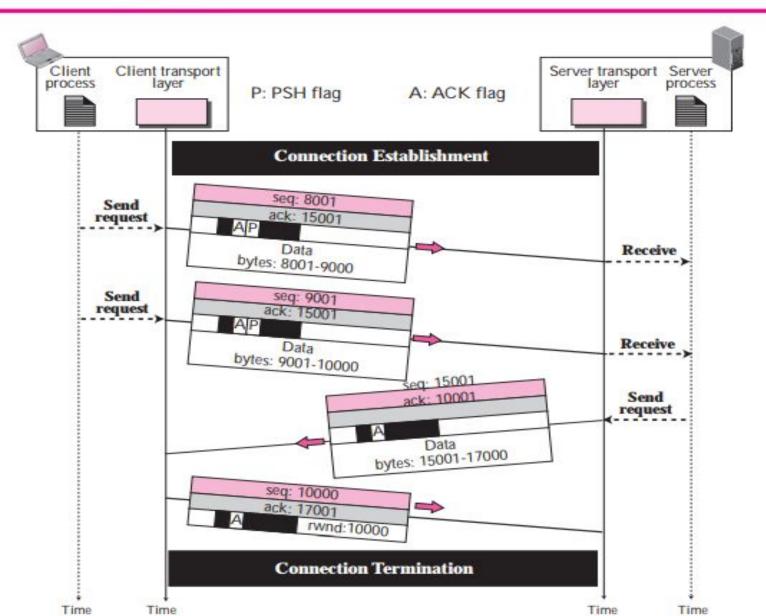
Note: SYN Flooding Attack TCP is vulnerable to SYN Flooding Attack The attacker sends repeated SYN packets to every port on the targeted server, often using a fake IP address. • The server, unaware of the attack, receives multiple, apparently legitimate requests to establish communication. • It responds to each attempt with a SYN-ACK packet from each open port. • The malicious client either does not send the expected ACK, or—if the IP address is spoofed—never receives the SYN-ACK Either way, the server under attack will wait for acknowledgement of its SYN-ACK packet for some time.

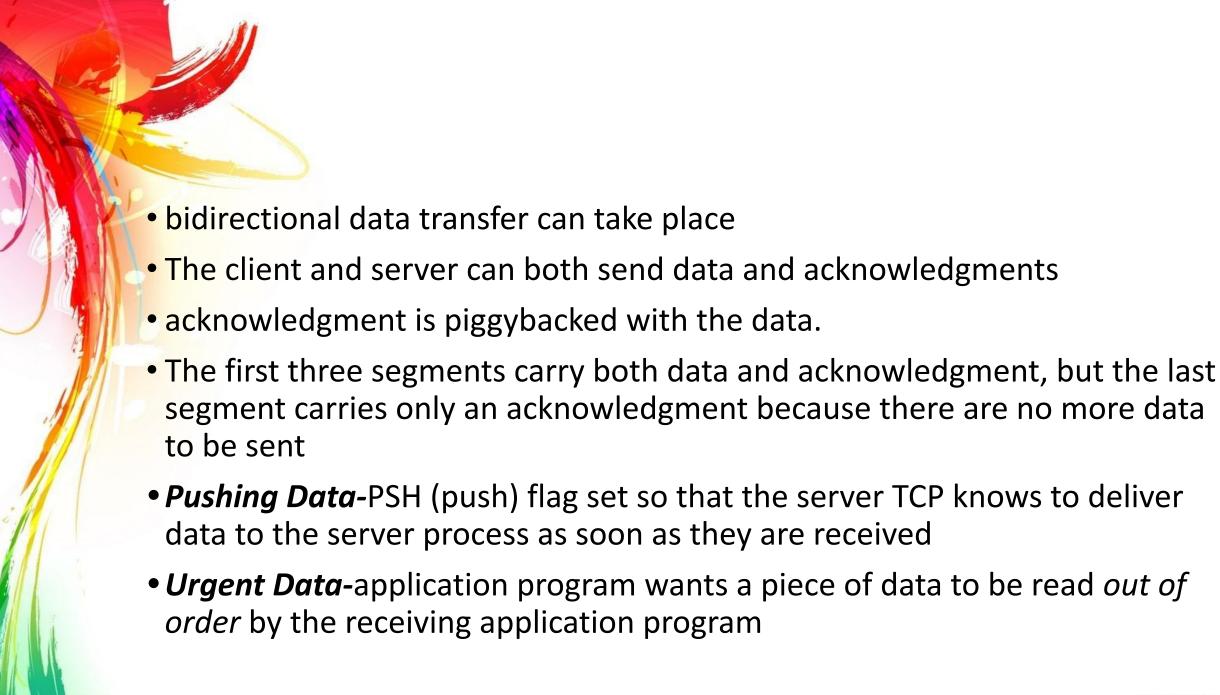


achieve network saturation.

Step 2: Data exchange in both direction

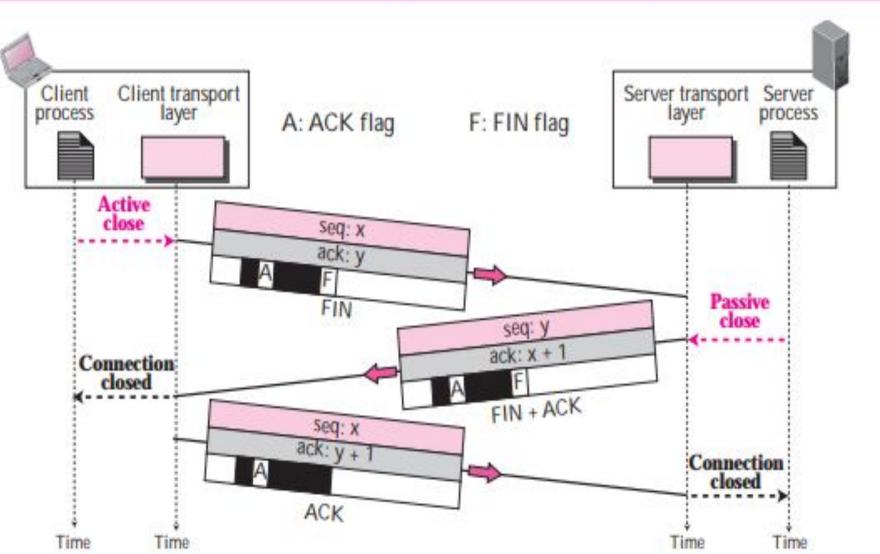
Data transfer





Step 3: Connection Termination

Connection termination using three-way handshaking

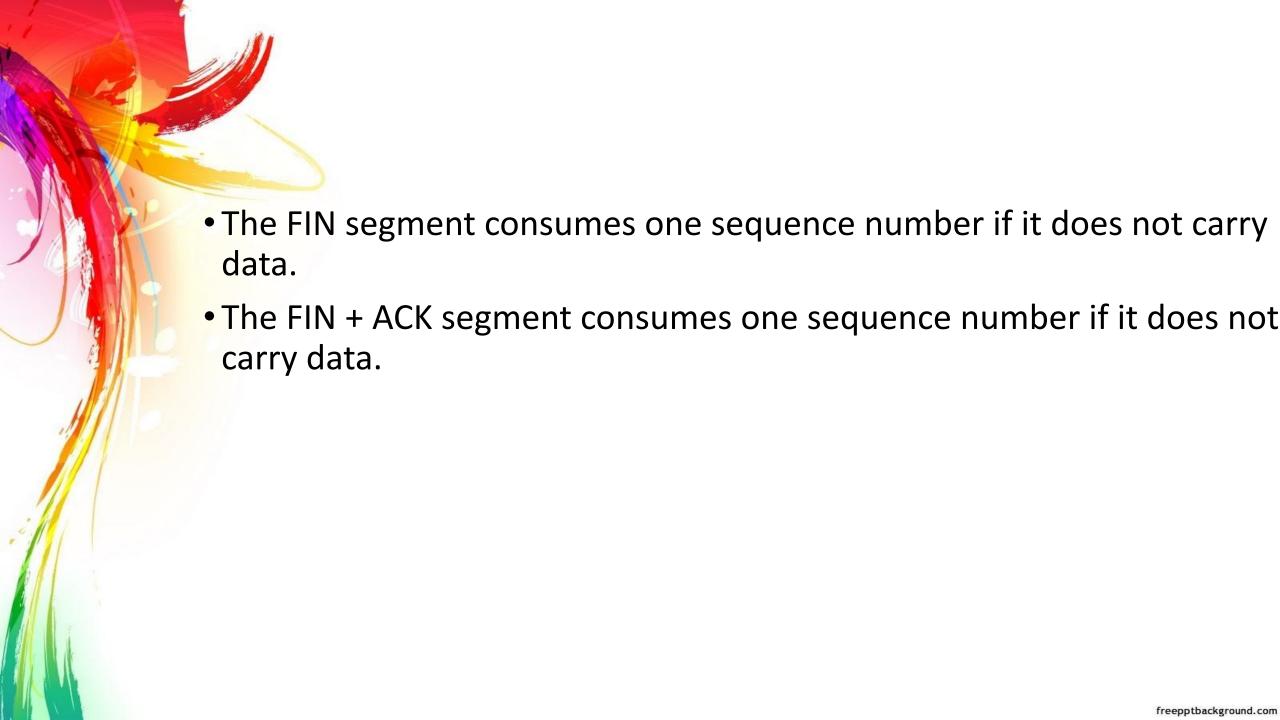


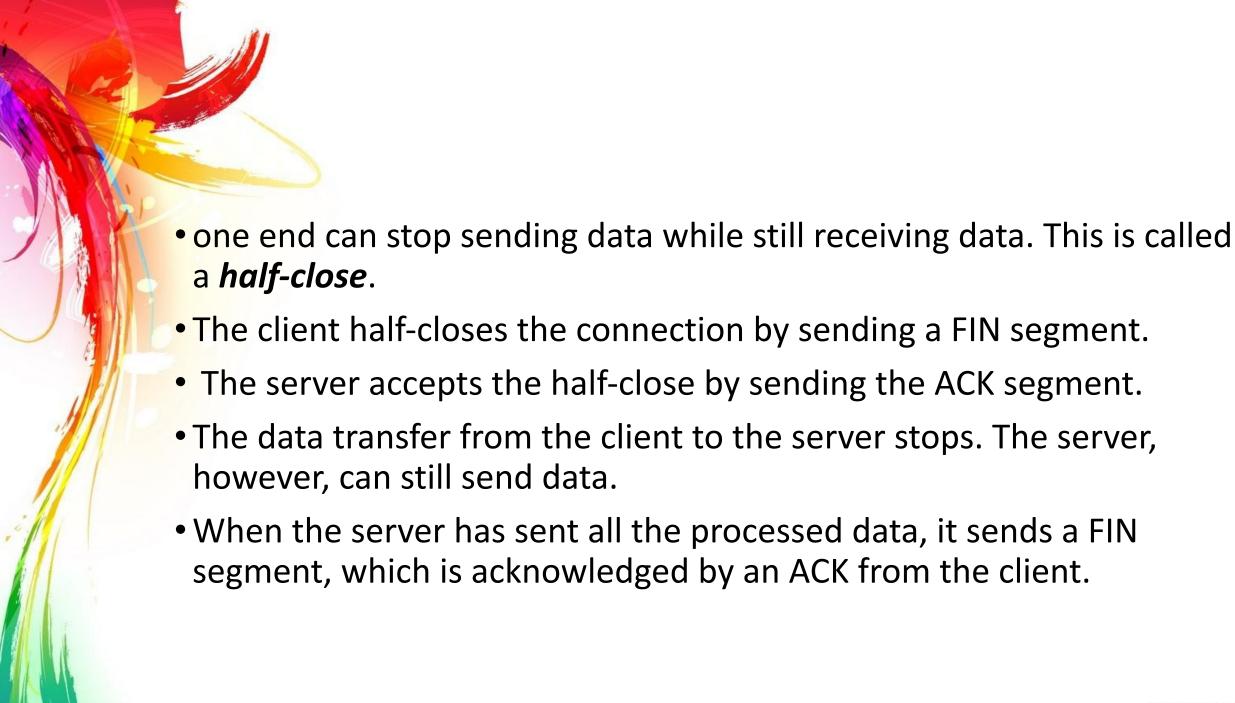


Two options for connection termination:

three-way handshaking

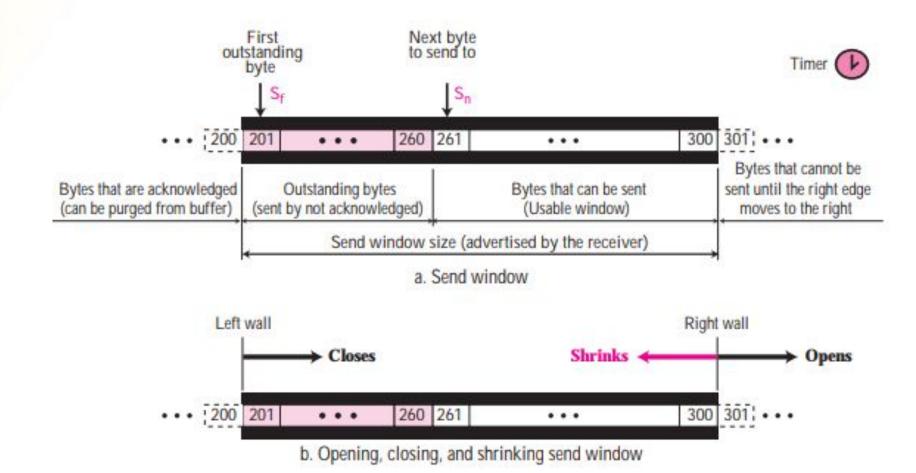
- To close a FIN segment in which the FIN flag is set
- FIN + ACK segment, to confirm the receipt of the FIN segment from the client and at the same time to announce the closing of the connection in the other direction
- last segment, an ACK segment, to confirm the receipt of the FIN segment from the TCP server
- four-way handshaking with a half-close option.





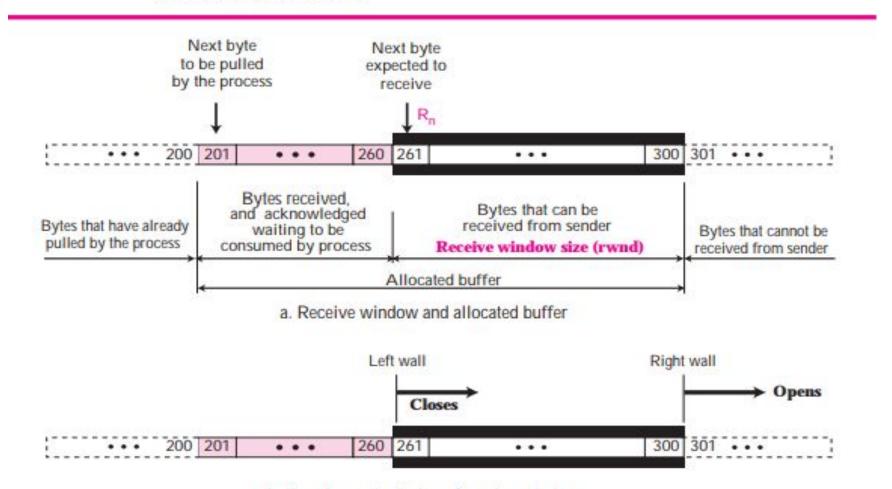
Windows in TCP for Flow Control

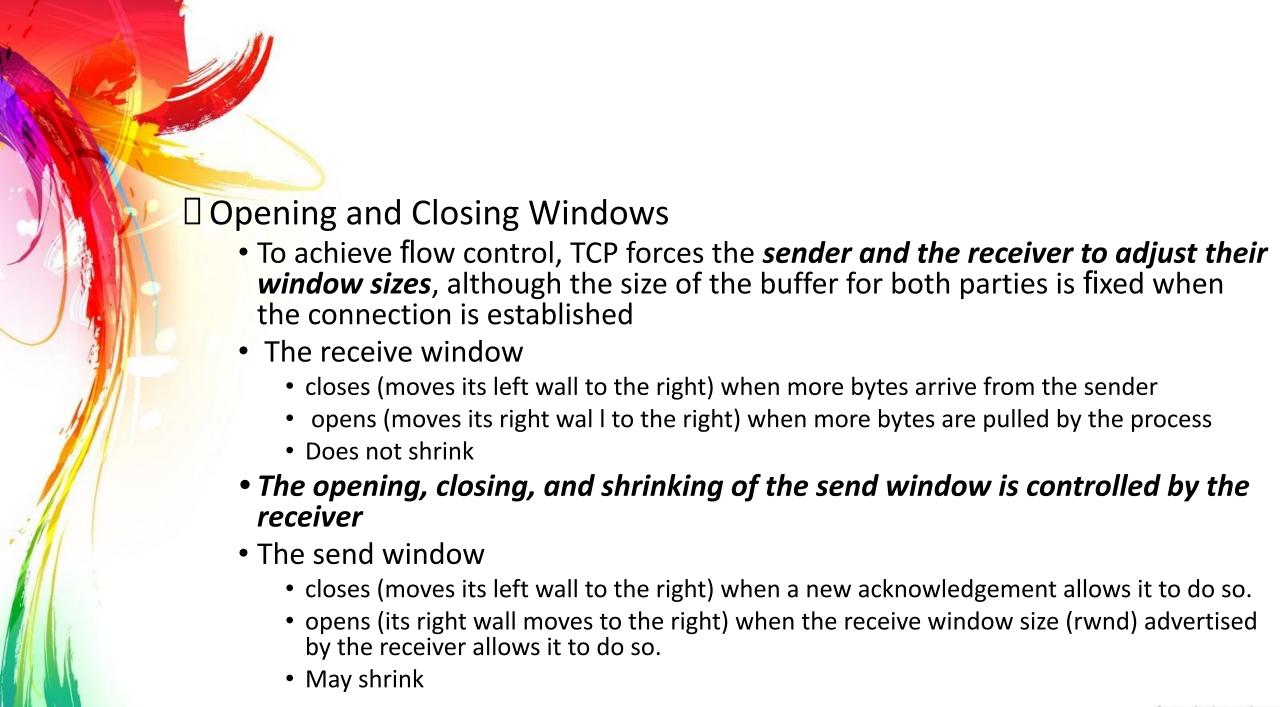
- Send Window
 - send window size is dictated by the receiver (flow control) and the congestion in the underlying network (congestion control).

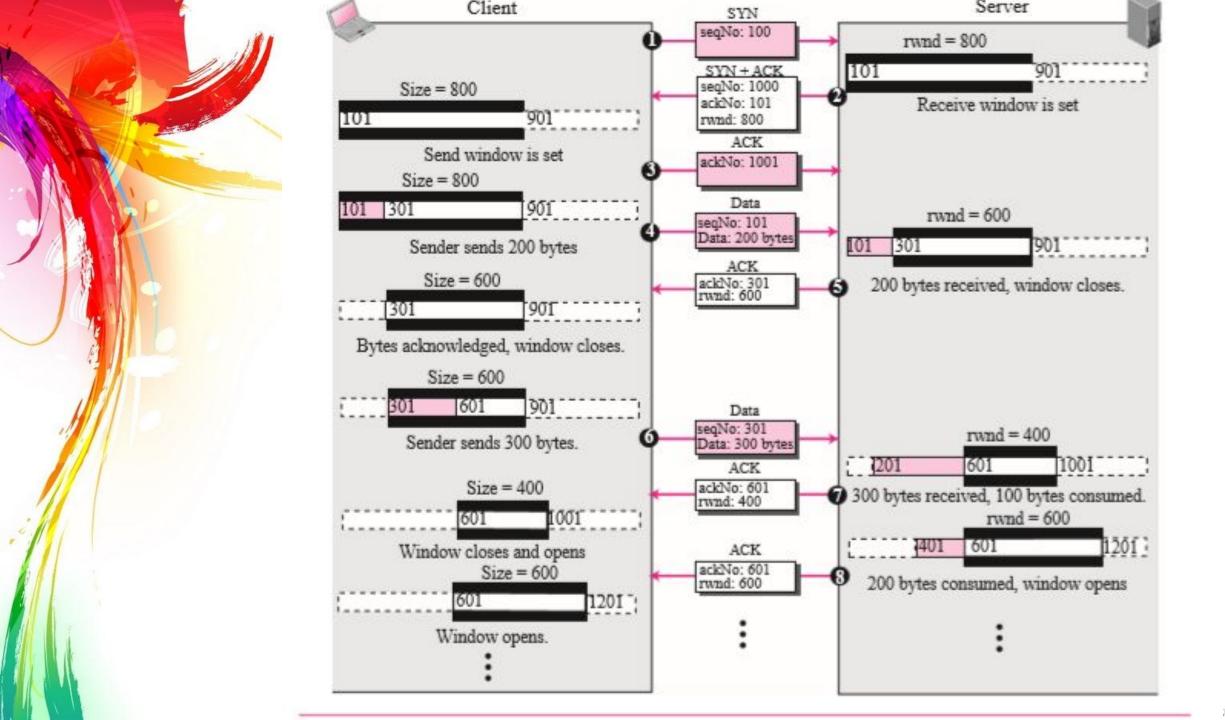


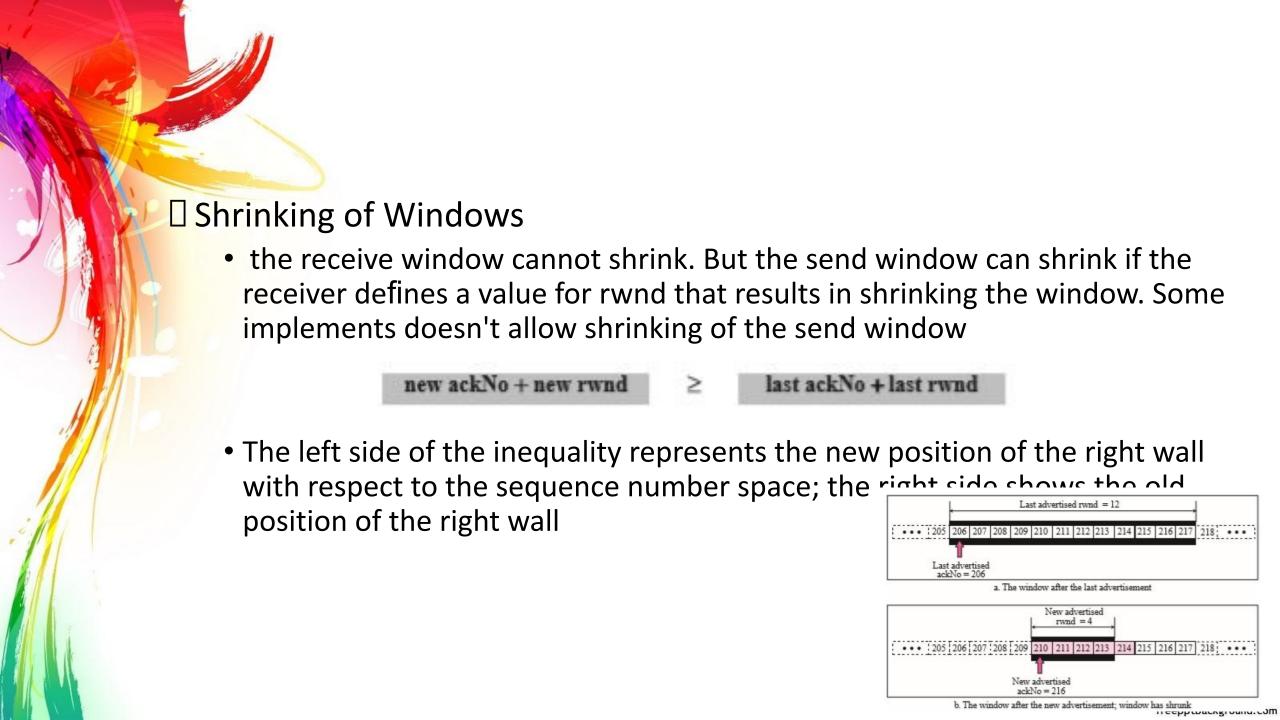


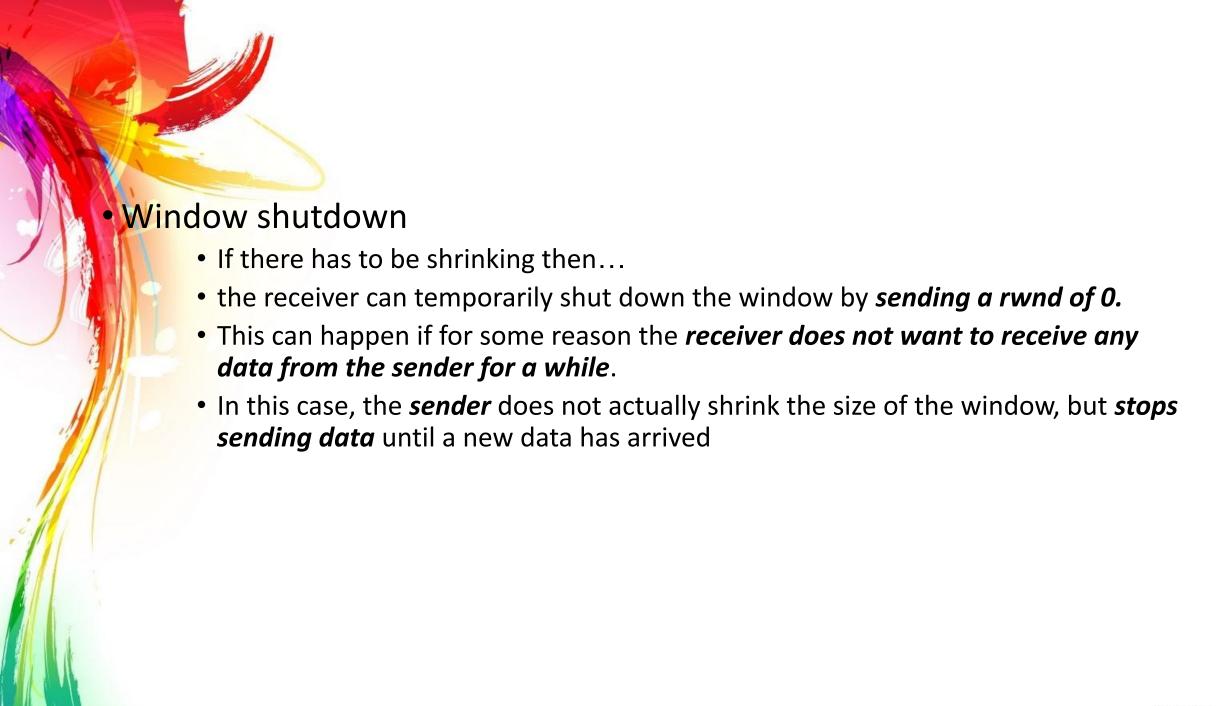
- Does not shrink like senders
- The receiver window size determines the number of bytes that the receive window can accept from the sender Receive window in TCP

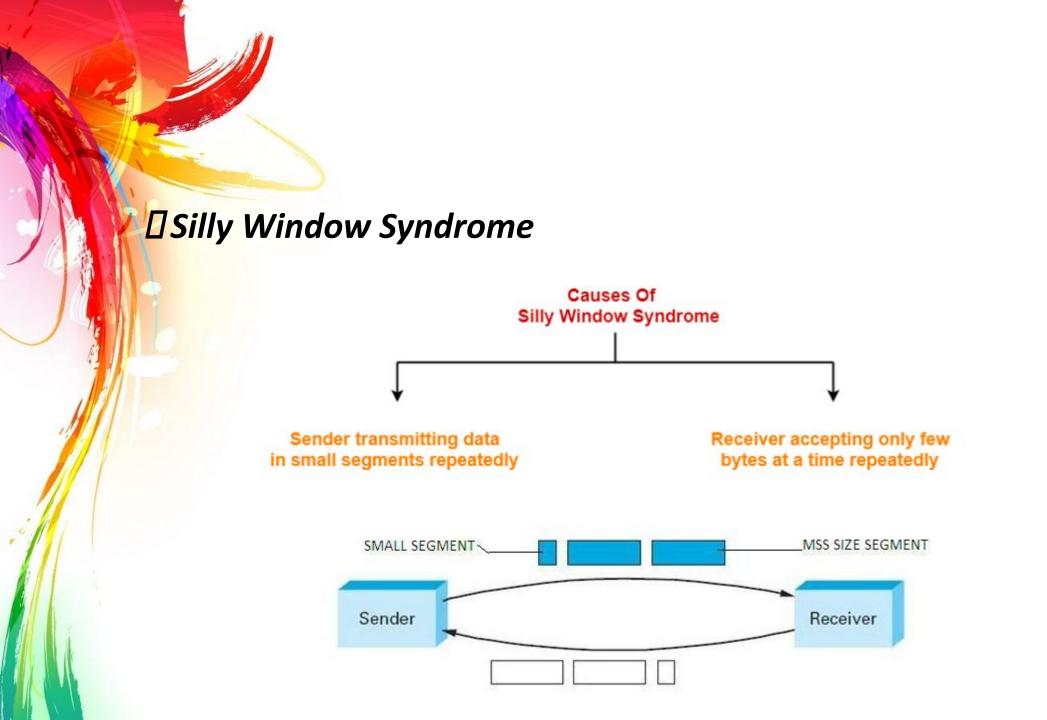


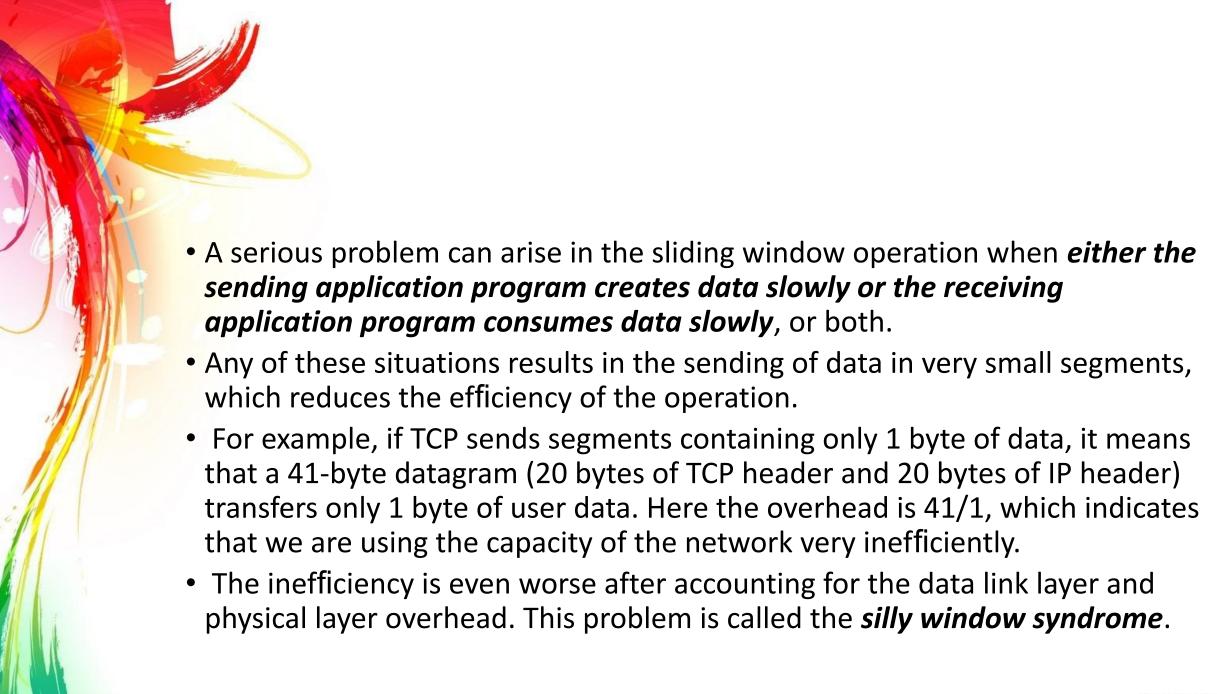












Syndrome Created by the Sender

- if it is serving an application program that creates data slowly, for example, 1 byte at a time.
- The application program writes 1 byte at a time into the buffer of the sending TCP.
- If the sending TCP does not have any specific instructions, it may create segments containing 1 byte of data. The result is a lot of 41-byte segments that are sent.
- How to prevent the sending TCP from sending the data byte by byte.
 Solution: The sending TCP must be forced to wait and collect data to send in a larger block.
- Problem ????

How long should the sending TCP wait? If it waits too long, it may delay the process. If it does not wait long enough, it may end up sending small segments Solution: **Nagle's algorithm**

Nagle's Algorithm

- 1. Sending TCP sends the first piece of data it receives from the sending application program even if it is only 1 byte.
- 2. Later, the sending TCP accumulates data in the output buffer and waits until either the receiving TCP sends an acknowledgment or until enough data has accumulated to fill a maximum-size segment (MSS).

At this time, the sending TCP can send the segment.

3. Step 2 is repeated for the rest of the transmission.



- it takes into account the *speed of the application program* that creates the data and the *speed of the network* that transports the data.
- If the application program is faster than the network, the segments are larger (maximum-size segments). If the application program is slower than the network, the segments are smaller (less than the maximum segment size).
- Nagle's algorithm has the advantage of increasing the efficiency of the network, but it has the disadvantage of increasing the delay. However, disabling Nagle's algorithm increases the network congestion, which could cause overall throughput degradation.

Syndrome Created by the Receive

- Receiving program that consumes data slowly, like 1 byte at a time.
- Suppose that the sending application program creates data in blocks of 1 kilobyte, but the receiving application program consumes data 1 byte at a time. Also suppose that the input buffer of the receiving TCP is 4 kilobytes. The sender sends the first 4 kilobytes of data. The receiver stores it in its buffer. Now its buffer is full. It advertises a window size of zero, which means the sender should stop sending data. The receiving application reads the first byte of data from the input buffer of the receiving TCP. Now there is 1 byte of space in the incoming buffer. The receiving TCP announces a window size of 1 byte, which means that the sending TCP, which is eagerly waiting to send data, takes this acknowledgment as good news and sends a segment carrying only 1 byte of data. The procedure will continue. One byte of data is consumed and a segment carrying 1 byte of data is sent.

Problem??? an efficiency problem and the silly window syndrome.

Solution:

- Clark's Solution
- Delayed Acknowledgment

Clark's Solution

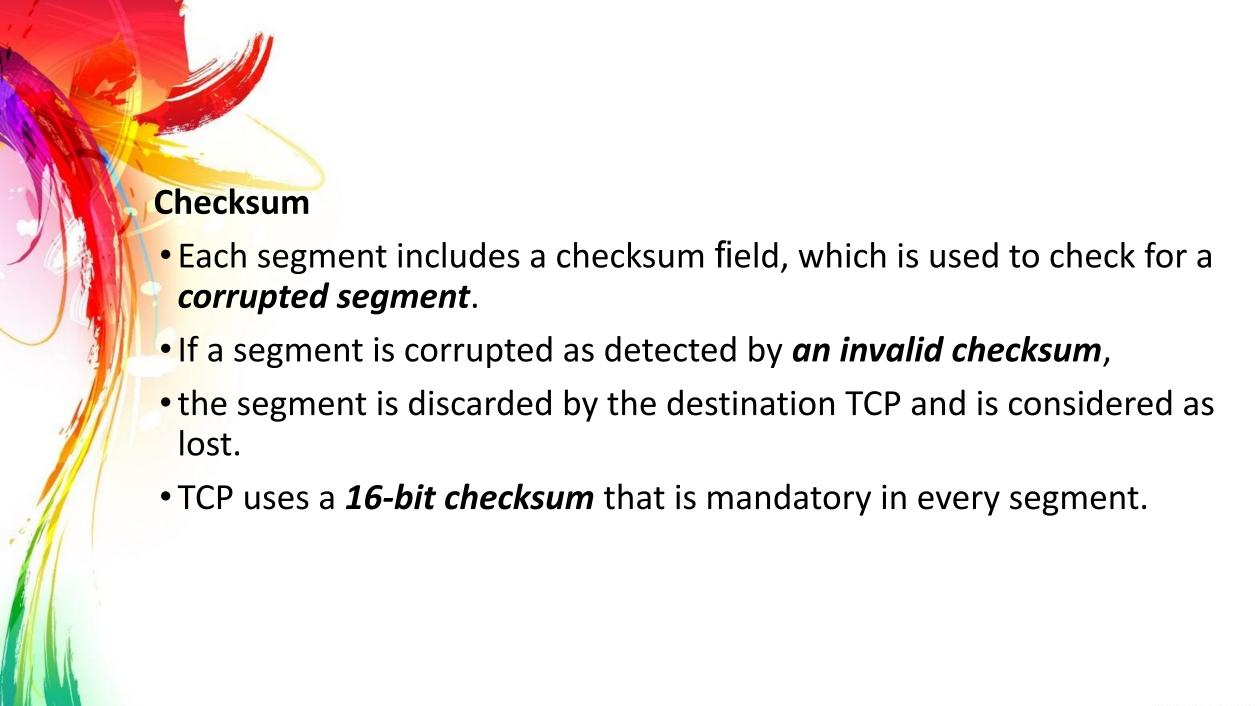
- **send an acknowledgment** as soon as the data arrive, but to announce a **window size of zero**
- until either there is enough space to accommodate a segment of maximum size or until at least half of the receive buffer is empty.

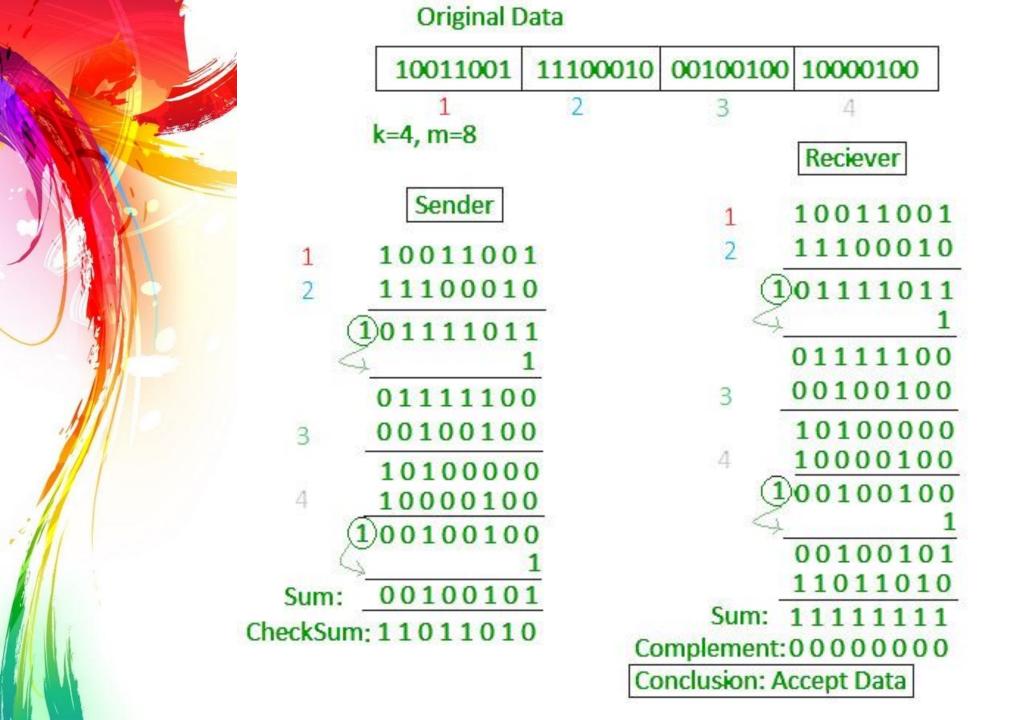
Delayed Acknowledgment

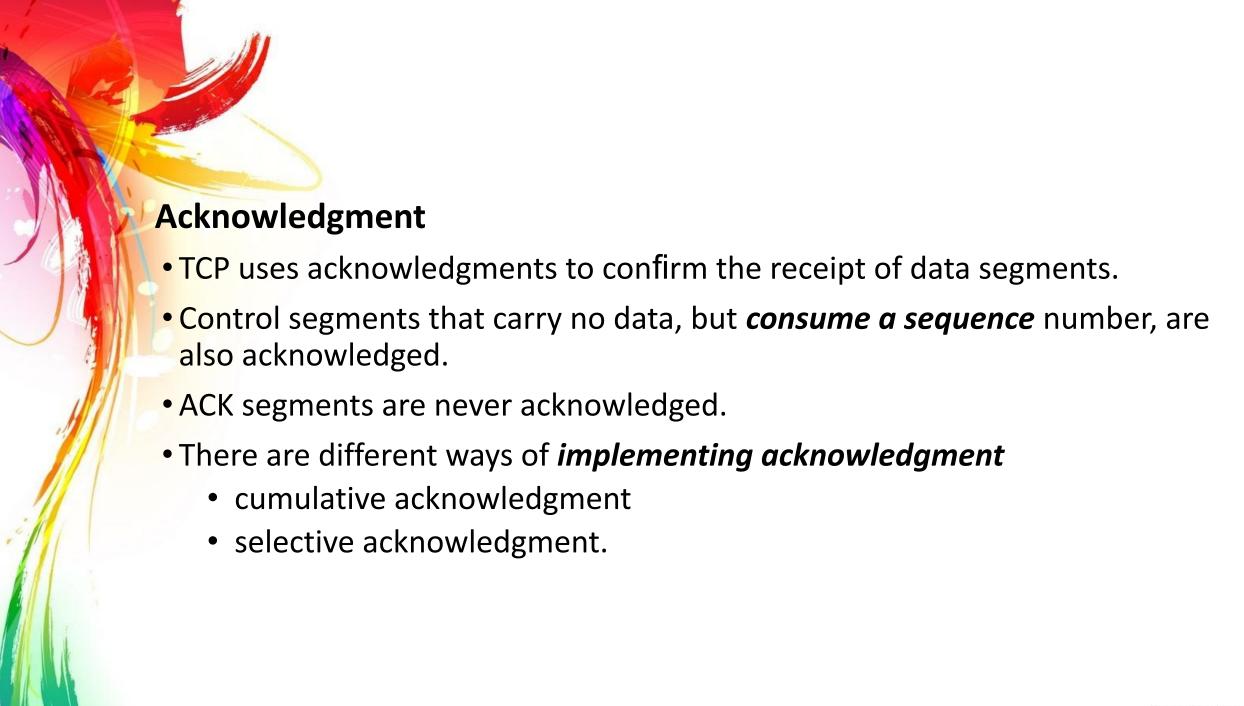
- delay sending the acknowledgment.
- Segment's arrives, it is not acknowledged immediately.
- Receiver waits until there is a decent amount of space in its incoming buffer before acknowledging the arrived segments.
- Advantage: it reduces traffic. The receiver does not have to acknowledge each segment
- Disadvantage: delayed acknowledgment may result in the sender unnecessarily retransmitting the unacknowledged segments
- Timer can be set to balance this situation

ERROR CONTROL

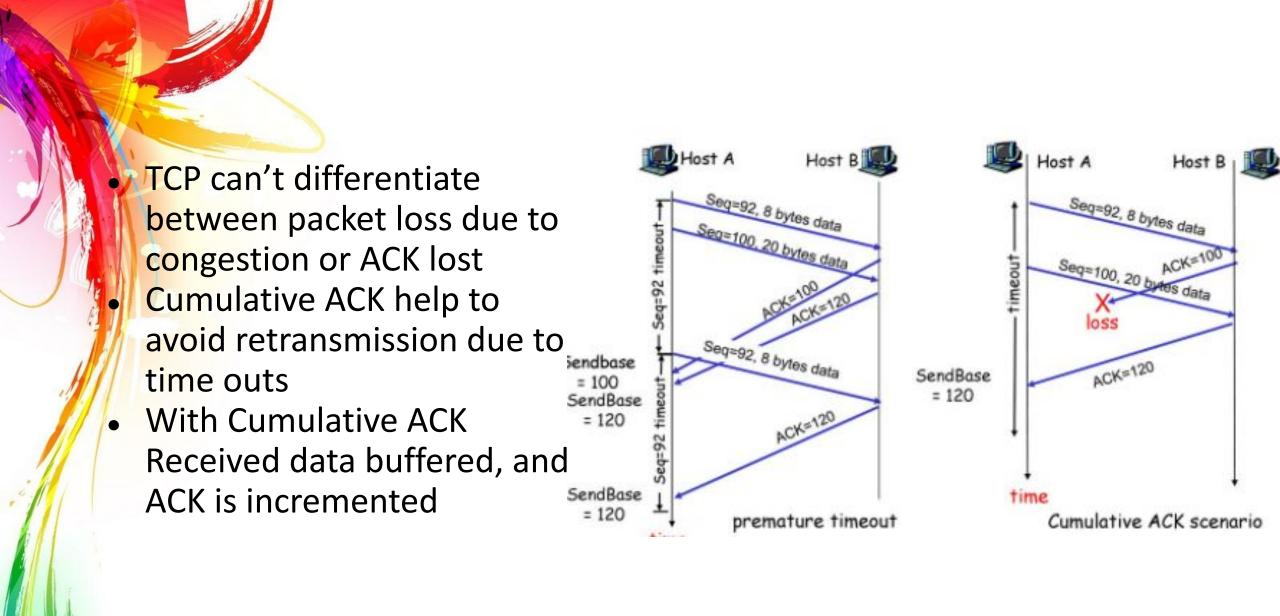
- TCP is a reliable transport layer protocol
- TCP provides reliability using error control
- Error control includes mechanisms for
 - detecting and resending corrupted segments
 - resending *lost segments*
 - storing out-of-order segments until missing segments arrive,
 - detecting and discarding duplicated segments.
- Three simple tools:
 - Checksum
 - Acknowledgment
 - Retransmission





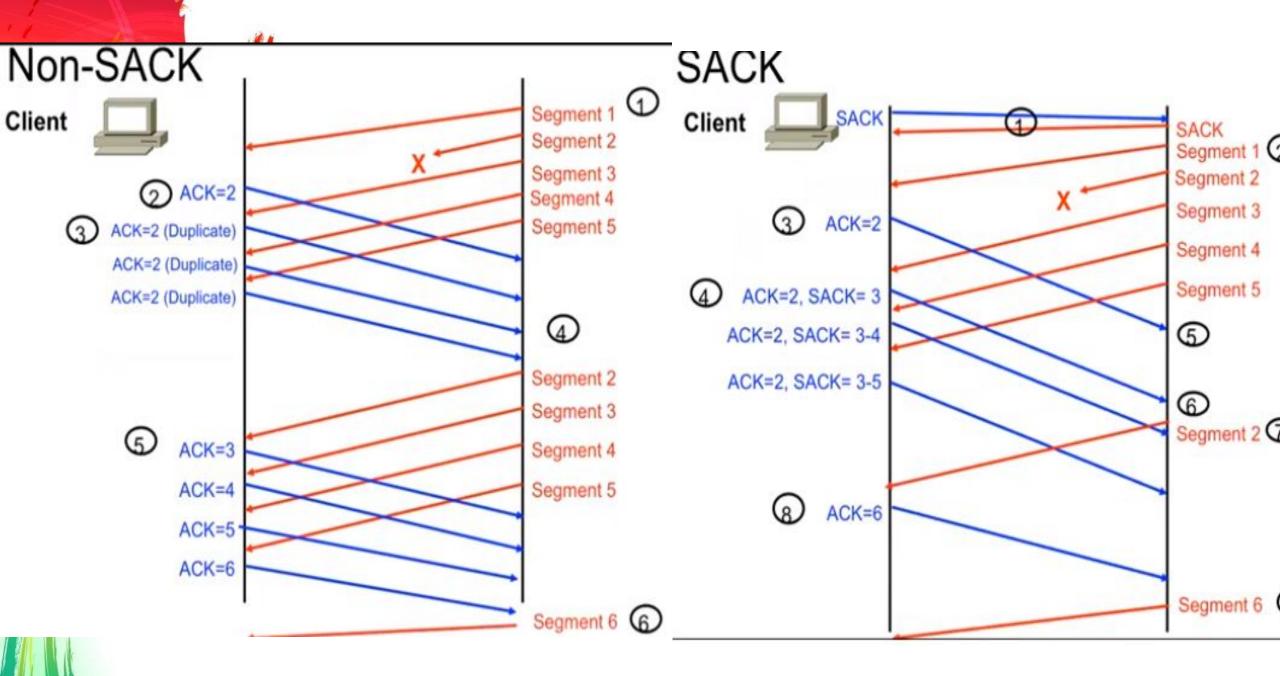


1. Cumulative Acknowledgment (ACK) TCP was originally designed to acknowledge receipt of segments cumulatively. The receiver advertises the next byte it expects to receive, ignoring all segments received and stored out of order. Also referred to as positive cumulative acknowledgment or ACK. The word "positive" indicates that no feedback is provided for discarded, lost, or duplicate segments. • The **32-bit** ACK field in the TCP header is used for cumulative acknowledgments and its value is valid only when the ACK flag bit is set to



2. Selective Acknowledgment (SACK)

- SACK does not replace ACK
- Reports additional information to the sender.
- A SACK reports a block of data that is out of order, and also a block of segments that is duplicated
- Since no provision in the TCP header for adding this type of information, SACK is implemented as an option at the end of the TCP header.
- it is an optional feature and is negotiated during 3 way handshake

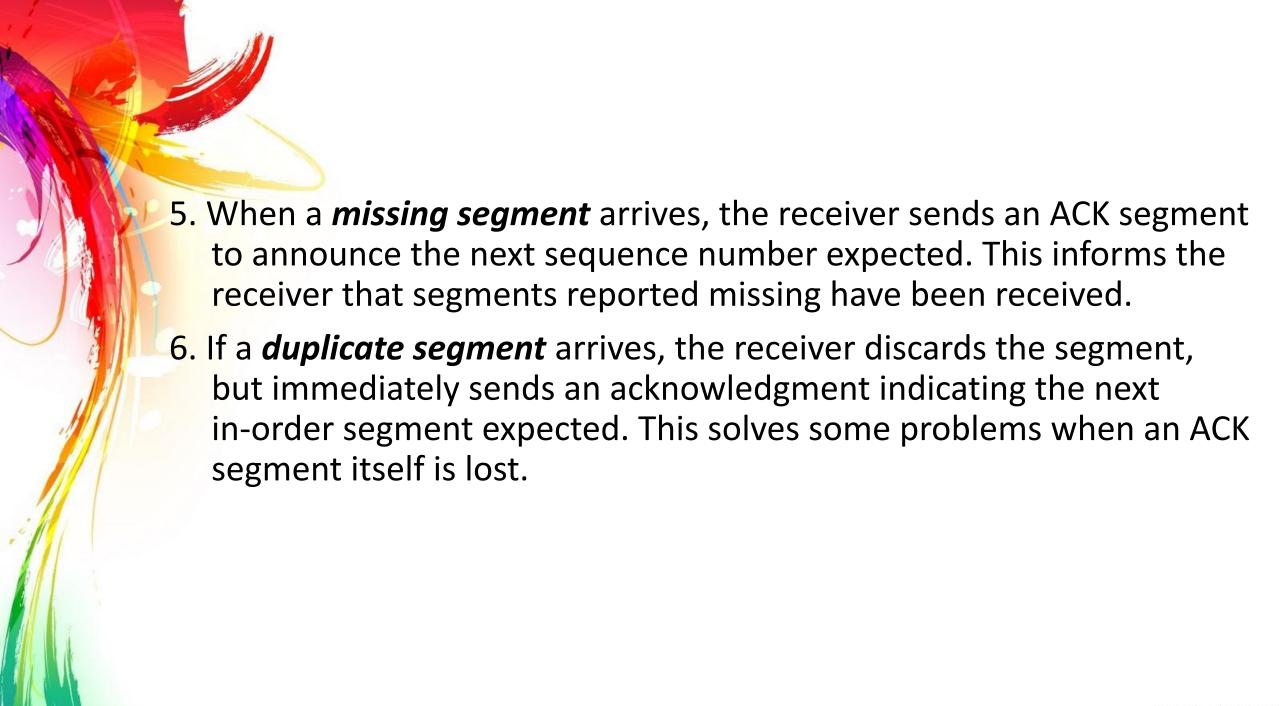


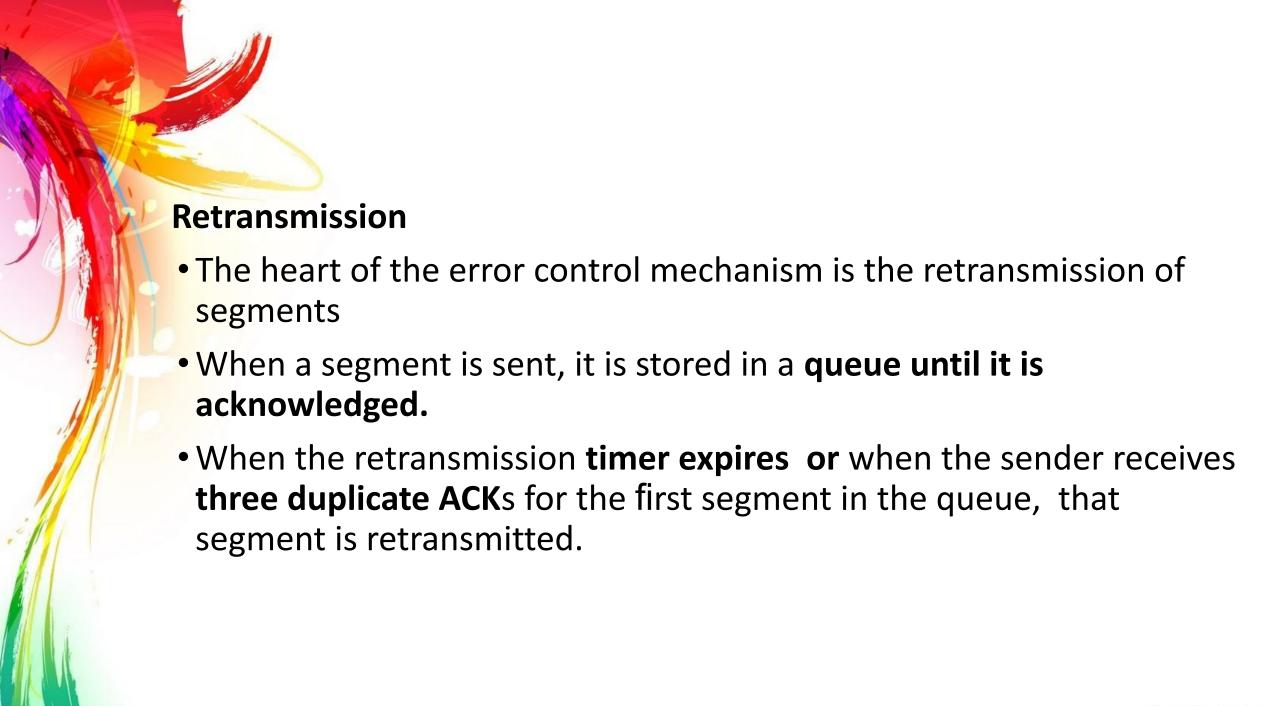
Summarizing how Acknowledgments are Generating

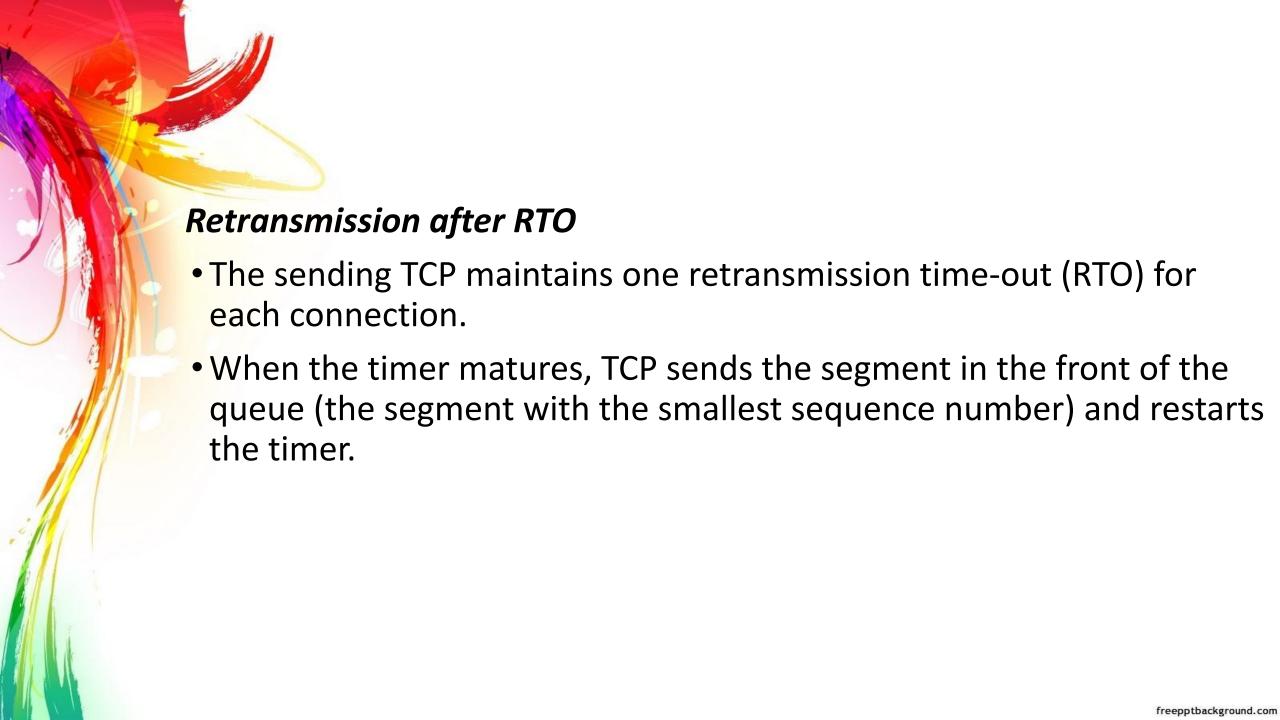
1. When end A sends a data segment to end B, it must include (*piggyback*) an acknowledgment that gives the next sequence number it expects to receive. This rule decreases the number of segments needed and therefore reduces traffic.

When the receiver has no data to send and it receives an in-order segment (with expected sequence number) and the previous segment has already been acknowledged, the receiver delays sending an ACK segment until another segment arrives or until a period of time (normally 500 ms) has passed. In other words, the receiver needs to *delay sending an ACK segment if there is only one outstanding in-order segment*. This rule reduces ACK segment traffic.

- 3. When a segment arrives with a sequence number that is expected by the receiver, and the previous in-order segment has not been acknowledged, the receiver immediately sends an ACK segment. In other words, there should not be more than two in-order unacknowledged segments at any time. This prevents the unnecessary retransmission of segments that may create congestion in the network.
 - 4. When a segment arrives with an *out-of-order sequence number* that is higher than expected, the receiver immediately sends an ACK segment announcing the sequence number of the next expected segment. This leads to the fast *retransmission of missing segments*.







Retransmission after Three Duplicate ACK Segments

- The previous rule about retransmission of a segment is sufficient if the value of RTO is not large.
- To improve throughput by allowing sender to retransmit sooner than waiting for a time out, most implementations today follow the three duplicate ACKs rule and retransmit the missing segment immediately.
- Also referred to as fast retransmission or Reno.
- In this version, if three duplicate acknowledgments (i.e., an original ACK plus three exactly identical copies) arrives for a segment, the next segment is retransmitted without waiting for the time-out.

CONGESTION CONTROL

- When too many packets are driven on the same link
- The queue overflows
- Packets get dropped

Network is congested!

Network should provide a congestion control mechanism to deal with such a situation, TCP uses a

- congestion window
- congestion policy (Control)
 - that avoid congestion and detect and reduce the after effects of congestion

- Congestion control and Resource Allocation
 - Two sides of the same coin
- If the network takes active role in allocating resources
 - The congestion may be avoided
 - No need for congestion control
- Allocating resources with any precision is difficult
 - Resources are distributed throughout the network
- Then recover from the congestion when it occurs
 - Easier approach but it can be disruptive because many packets many be discarded by the network before congestions can be controlled

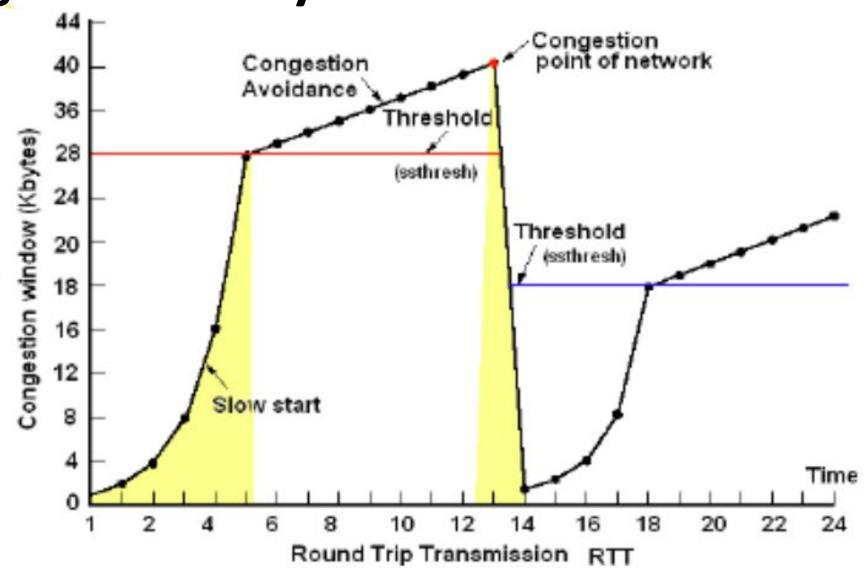


Congestion Window

- Only the receiver can dictate to the sender the size of the sender's window! Not really
- Another entity that is affected is the network.
- If the network cannot deliver the data as fast as it is created by the sender, it must tell the sender to slow down.
- In addition to the receiver, the network is a second entity that determines the size of the sender's window.
- The sender has two pieces of information:
 - receiver-advertised window size and
 - congestion window size.
- The actual size of the window is the minimum of these two.

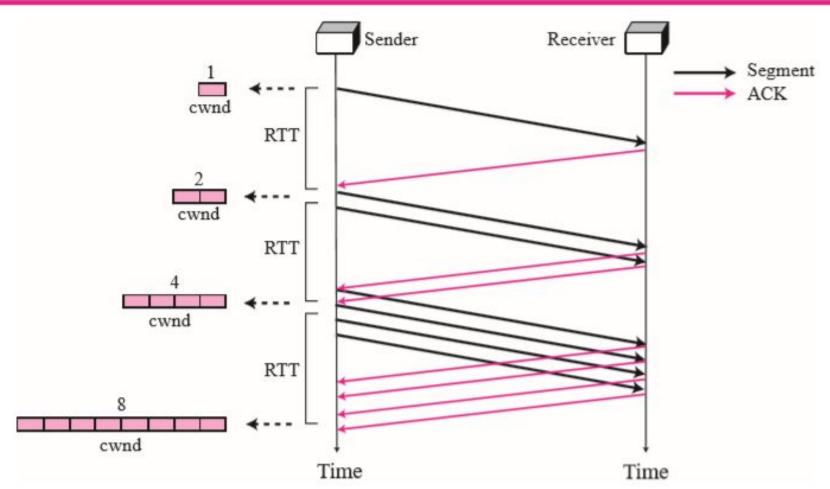
Actual window size = minimum (rwnd, cwnd)

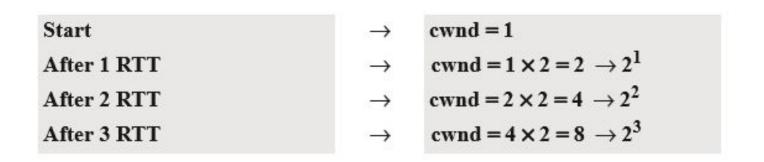
Congestion Policy



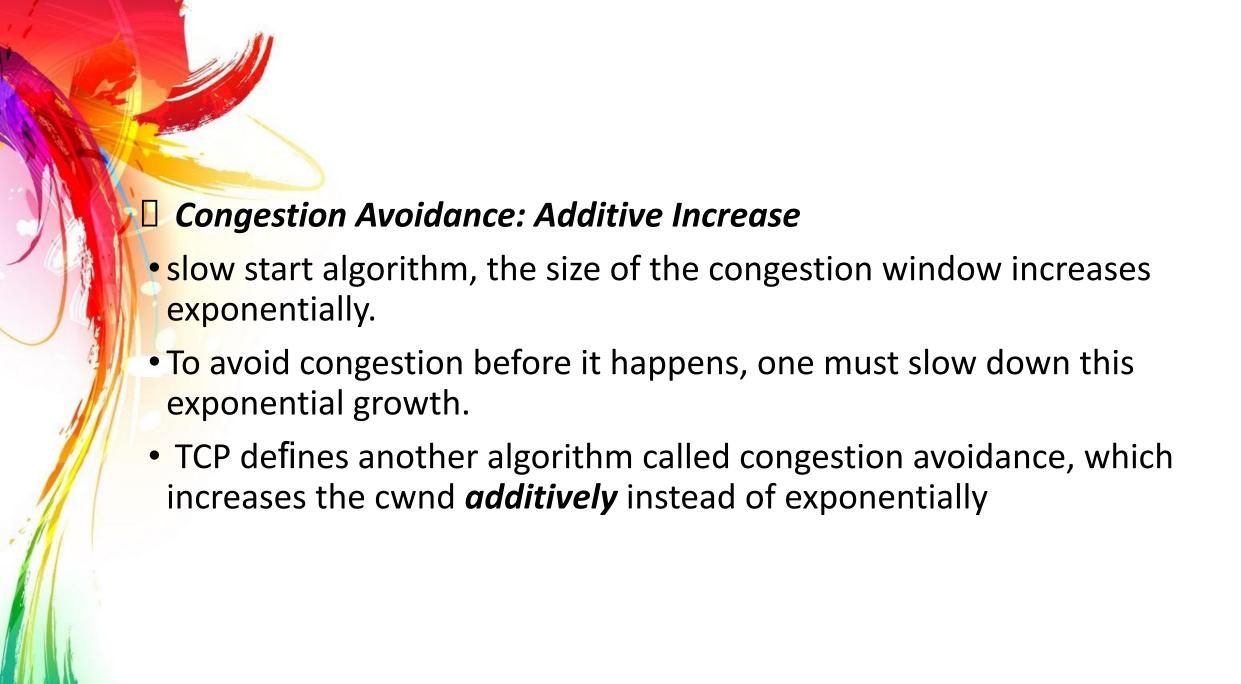
Congestion Policy Three phases: Slow start congestion avoidance congestion detection. ☐ Slow Start: Exponential Increase In the slow start phase, the sender starts with a slow rate of transmission, but increases the rate rapidly to reach a threshold. When the threshold is reached, the rate is reduced. Finally if ever congestion is detected, the sender goes back to the slow start or congestion avoidance phase

Slow start, exponential increase





- starts with one maximum segment size (MSS)
- size of the window increases one MSS each time one acknowledgement arrives
- Slow start cannot continue indefinitely.
- There is a threshold to stop this phase.
- The sender keeps track of a variable named *ssthresh* (slow start threshold).
- When the size of window in bytes reaches this threshold, slow start stops and the next phase starts.



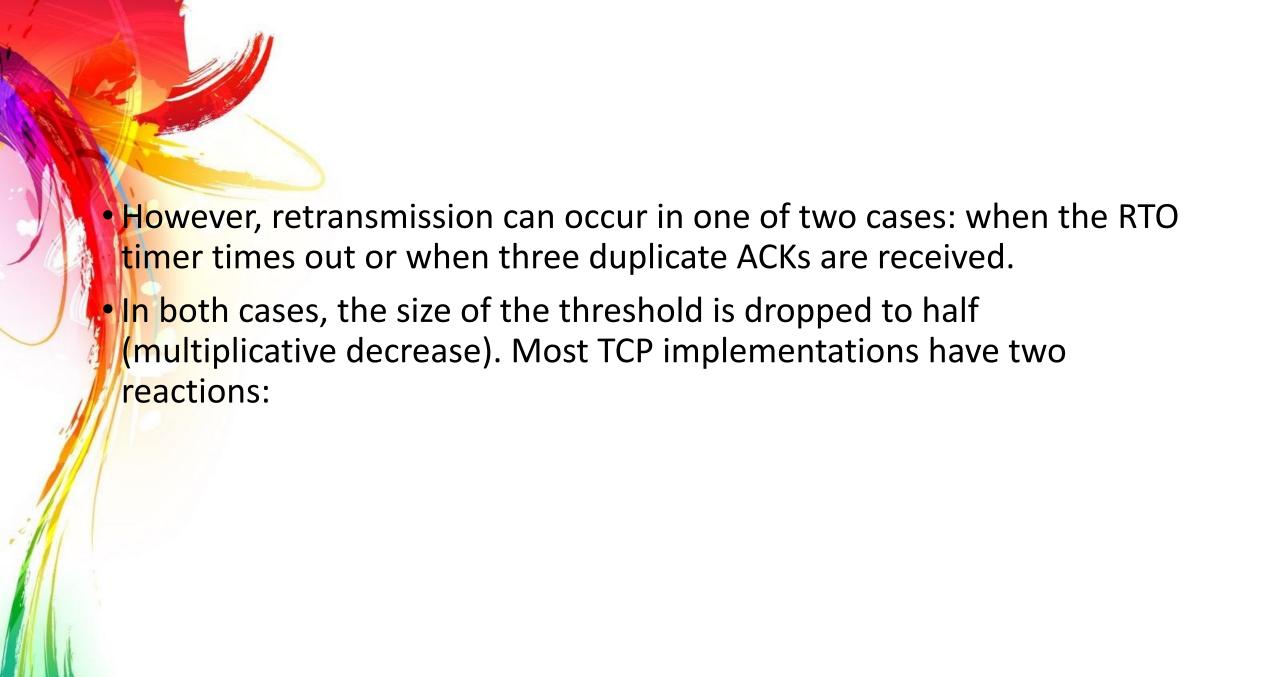
Congestion avoidance, additive increase Sender Receiver i = 4▶ Segment▶ ACK cwnd RTT i+1cwnd RTT i+2cwnd RTT i+3cwnd Time Time

- each time the whole "window" of segments is acknowledged, the size of the congestion window is increased by one.
- A window is the number of segments transmitted during RTT.
- The increase is based on RTT, not on the number of arrived ACKs as in slow start
- Start with slow start, after the sender has received acknowledgments for a complete window-size of segments, the size of the window is increased one segment

```
Start\rightarrow\operatorname{cwnd} = iAfter 1 RTT\rightarrow\operatorname{cwnd} = i + 1After 2 RTT\rightarrow\operatorname{cwnd} = i + 2After 3 RTT\rightarrow\operatorname{cwnd} = i + 3
```

Congestion Detection: Multiplicative Decrease

- If congestion occurs, the congestion window size must be decreased.
- The only way a sender can guess that congestion has occurred is the need to retransmit a segment. This is a major assumption made by TCP.
- Retransmission is needed to recover a missing packet which is assumed to have been dropped (i.e., lost) by a router that had so many incoming packets, that had to drop the missing segment, i.e., the router/network became overloaded or congested.



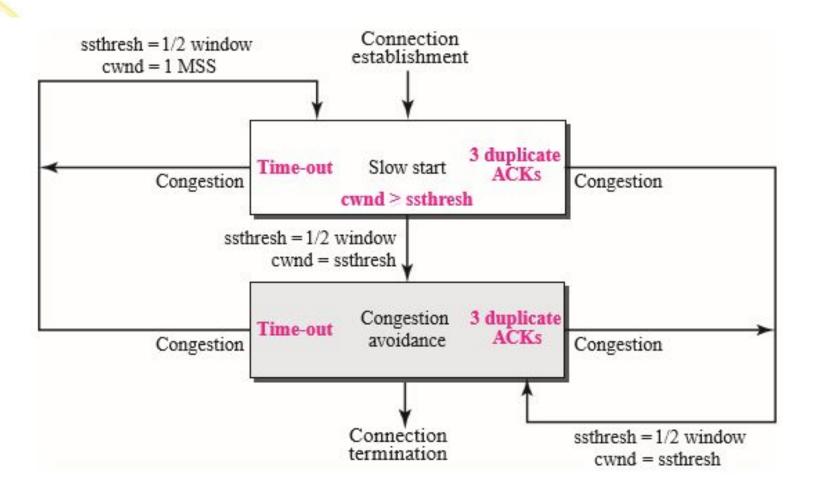
1. If a time-out occurs, there is a stronger possibility of congestion; a segment has probably been dropped in the network and there is no news about the following sent segments. In this case TCP reacts strongly: It sets the value of the threshold to half of the current window size. b. It reduces cwnd back to one segment. c. It starts the slow start phase again.

2. If three duplicate ACKs are received, there is a weaker possibility of congestion; a segment may have been dropped but some segments after that have arrived safely since three duplicate ACKs are received. This is called fast transmission and fast recovery.

In this case, TCP has a weaker reaction as shown below:

- a. It sets the value of the threshold to half of the current window size.
- b. It sets cwnd to the value of the threshold (some implementations add three segment sizes to the threshold).
- c. It starts the congestion avoidance phase.

TCP Congestion Policy Summary



TCP Timers

- 4 types of timers
 - TIME-WAIT Timer
 - Retransmission Timer
 - Persistence Time
 - Keepalive Timer

TIME WAIT Timer

- > This timer during connection termination
- > The timer starts after sending the last Ack for 2nd FIN and closing the connection.

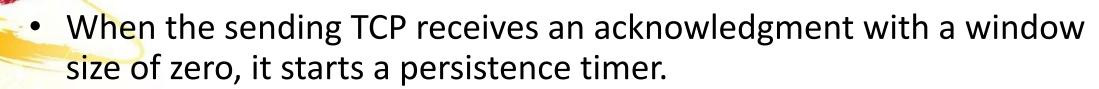
Retransmission Timer

To retransmit lost segments, TCP employs one retransmission timer (for the whole connection period) that handles the retransmission time-out (RTO), the waiting time for an acknowledgment of a segment.

Persistence Timer

- To deal with a zero-window-size advertisement, TCP needs another timer.
- If the receiving TCP announces a window size of zero, the sending TCP stops transmitting segments until the receiving TCP sends an ACK segment announcing a nonzero window size.
- This ACK segment can be lost. (ACK segments are not acknowledged nor retransmitted in TCP)

- If this acknowledgment is lost, the receiving TCP thinks that it has done its job and waits for the sending TCP to send more segments.
- There is no retransmission timer for a segment containing only an acknowledgment.
- The sending TCP has not received an acknowledgment and waits for the other TCP to send an acknowledgment advertising the size of the window.
- Both TCPs might continue to wait for each other forever (a deadlock).
- To correct this deadlock, TCP uses a persistence timer for each connection.



 When the persistence timer goes off, the sending TCP sends a special segment called a probe.

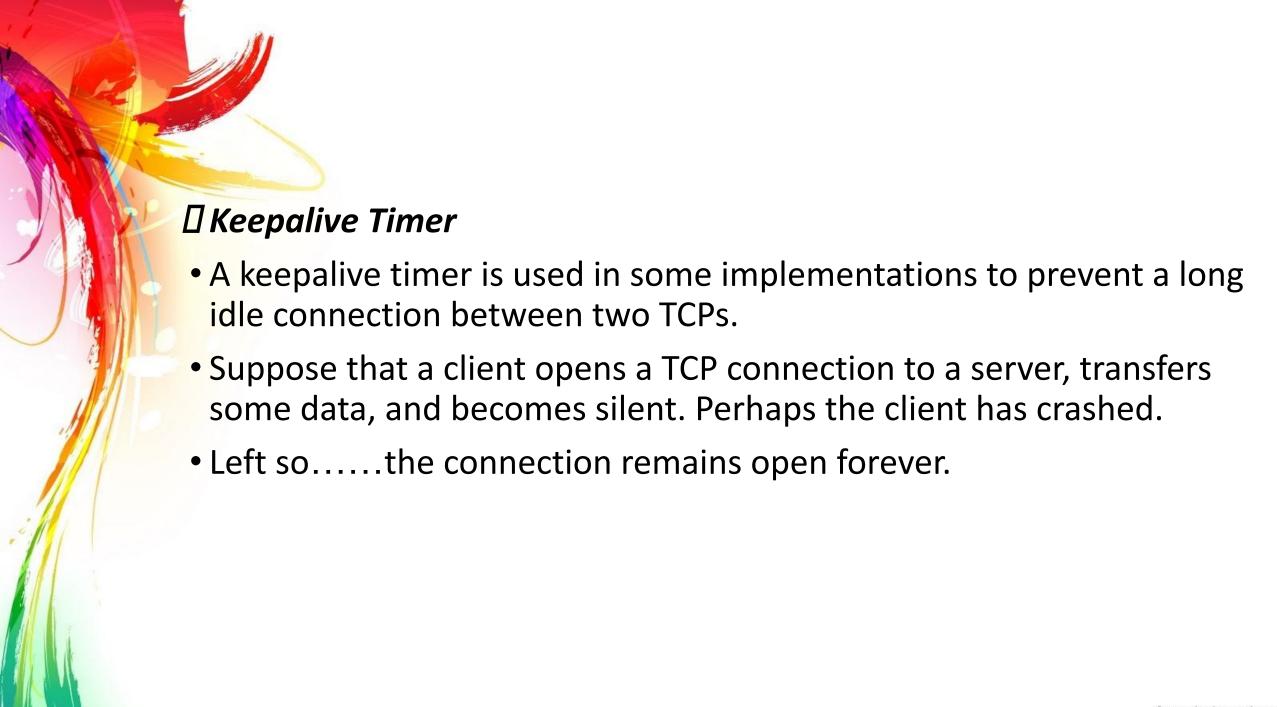
This segment contains only 1 byte of new data.

It has a sequence number, but its sequence number is never acknowledged

The probe causes the receiving TCP to resend the acknowledgment.

• The value of the persistence timer is set to the value of the retransmission time.

- if a response is not received from the receiver, another probe segment is sent and the value of the persistence timer is doubled and reset.
- The sender continues sending the probe segments and doubling and resetting the value of the persistence timer until the value reaches a threshold
- After that the sender sends one probe segment every 60 s until the window is reopened.





- Each time the server hears from a client, it resets this timer.
- The time-out is usually 2 hours. If the server does not hear from the client after 2 hours, it sends a probe segment.
- If there is no response after 10 probes, each of which is 75 s apart, it assumes that the client is down and terminates the connection.